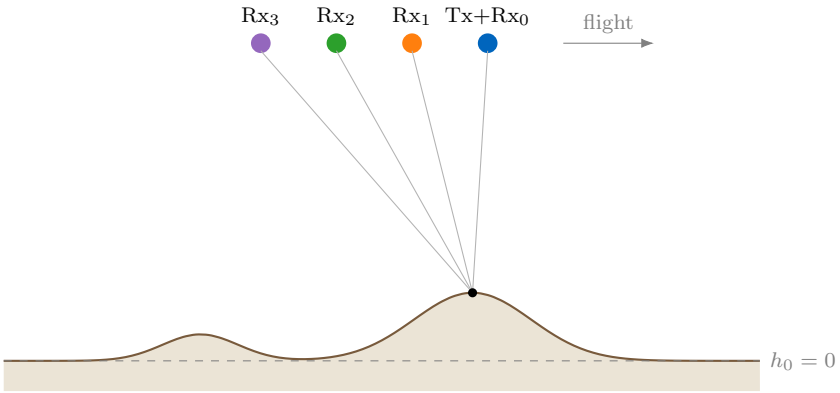
DLR-DAAD Research Fellowship

Microwaves and Radar Institute (DLR-HR), Oberpfaffenhofen

Topography-Dependent Azimuth Reconstruction in Multistatic SAR Systems

Project report



Lucas Ayub Rodrigues Tuler

Instituto Tecnológico de Aeronáutica (ITA), Brazil

Draft — October 2026

Abstract

Multichannel and multistatic synthetic aperture radar (SAR) systems for high-resolution wide-swath (HRWS) imaging sample the azimuth signal of each receiver at a pulse repetition frequency (PRF) below the Doppler bandwidth and recover the unambiguous spectrum by combining the channels in an azimuth reconstruction. The reconstruction filters are derived for a reference geometry. When the receivers are separated by cross-track baselines, a scatterer at a height different from the reference height carries a channel-dependent phase residual, the reconstruction no longer cancels the aliased parts of the spectrum, and azimuth ambiguities appear.

This report studies that topography-dependent mismatch with a numerical simulator of a spaceborne formation ($\lambda = 0.25\,\mathrm{m}$, $r_0 = 766\,\mathrm{km}$, $N_{\mathrm{rx}} = 2$ to 8 receivers, cross-track baselines up to 500 m) and compares correction strategies that use *only the system parameters and a digital elevation model (DEM)*, never the positions of the targets. The sub-aperture topography- and aperture-dependent (SATA) algorithm, applied to each channel before the reconstruction, is shown to be exact only when the terrain has the same height at every ground position folded onto the same Doppler bin, i.e. at positions $X \approx 6.2\,\mathrm{km}$ apart; on general terrain it leaves a structural residual that is predicted from the geometry within about 1 dB. A *joint* reconstruction that places the DEM phase of every folded parcel inside the inversion removes this limitation. Its final version models the channels from the focused line and the DEM and solves a regularised least-squares problem with conjugate gradients.

Over 1040 one-dimensional runs (52 scenes, $N_{\mathrm{rx}} = 2$–8) the median highest ambiguity drops from $-8.4\,\mathrm{dB}$ without correction to $-40.6\,\mathrm{dB}$ with the joint reconstruction (reference $-47.6\,\mathrm{dB}$). The two-dimensional extension, which applies the DEM at every pixel of a 3-D scene, reaches a median of $-44.4\,\mathrm{dB}$ over 640 runs (no correction $-13.9\,\mathrm{dB}$, reference $-47.6\,\mathrm{dB}$). The remaining failures are configurations whose baselines do not separate the folded parcels (two channels, equally spaced baselines) and are a property of the data rather than of the algorithm.

Contents

Abstract i

Symbols and acronyms vii

1 Introduction 1

1.1 Context . . . 1

1.2 Objective . . . 1

1.3 Contributions . . . 2

1.4 Outline . . . 2

2 Acquisition geometry and signal model 3

2.1 Formation and parameters . . . 3

2.2 Signal model . . . 3

2.3 Channel transfer functions and the coefficients C_0, C_1, C_2, Δt . . . 3

2.4 The topographic residual ΔC_0 . . . 7

3 Multichannel azimuth reconstruction and topography 9

3.1 Aliased channels and folded parcels . . . 9

3.2 How topography creates ambiguities . . . 9

3.3 Ambiguity metric . . . 12

4 SATA before the reconstruction 13

4.1 The SATA algorithm . . . 13

4.1.1 Why a sub-aperture is an image . . . 13

4.1.2 The steps . . . 13

4.2 SATA on each channel, then the standard reconstruction . . . 14

4.2.1 Whole band . . . 14

4.2.2 Per sub-band . . . 14

4.2.3 Why it is exact only on special terrain . . . 15

4.2.4 Verification . . . 16

4.2.5 SATA on simulated scenes . . . 17

4.2.6 Correction maps that use the target positions . . . 17

5 Joint topography correction and reconstruction 20

5.1 Joint reconstruction in the sub-apertures . . . 20
5.2 Joint reconstruction over the whole line . . . 20
5.3 Building the model and solving the system . . . 22
5.3.1 Building the model . . . 22
5.3.2 What is stored and what is applied . . . 25
5.3.3 The problem that is solved . . . 26
5.3.4 Conjugate gradients for least squares . . . 26
5.3.5 The answer: $U = K \odot \mathbf{F}\hat{\boldsymbol{g}}$ . . . 27
5.4 Properties . . . 28

6 Extension to two dimensions **30**
6.1 Model . . . 30
6.2 The common range band . . . 30

7 Results **34**
7.1 All methods with a DEM only . . . 34
7.2 Why the joint reconstruction is the final method . . . 35
7.3 Sweep of the final joint reconstruction in 1-D . . . 37
7.4 Sweep of the joint reconstruction in 2-D . . . 38
7.5 Limits . . . 40
7.6 Real terrain: SRTM lines over the Sierra Nevada . . . 40

8 Conclusions and outlook **48**
8.1 Conclusions . . . 48
8.2 Outlook . . . 48

List of Figures

2.1 Bistatic geometry of one channel 4
2.2 DPCA arrangement 4
2.3 Coefficients of the reference point 6
2.4 Bistatic minus monostatic path and its parabola 6
2.5 Change of the coefficients with the height of the scatterer 7
2.6 Origin of the topographic residual 8
2.7 Phase of the topographic residual 8

3.1 Folding of the Doppler spectrum 10
3.2 Aliased channel and reconstructed spectrum 10
3.3 Standard reconstruction of a target at the reference height and above it 11
3.4 Highest ambiguity of the standard reconstruction against height 11

4.1 The SATA kernel 14
4.2 The two SATA variants 15
4.3 The conflict of one bin 16
4.4 SATA on terrains with and without the period of the folds 18
4.5 Prediction of the level left by SATA 19
4.6 SATA can degrade a target at the reference height 19

5.1 Flowchart of the joint reconstruction in the sub-apertures 21
5.2 The model of the whole-line joint reconstruction 22
5.3 Flowchart of the joint reconstruction over the whole line 23
5.4 SATA + reconstruction versus the joint reconstruction 24
5.5 Coupling between Doppler bins in the joint model 26
5.6 From the channels to the image 28
5.7 Convergence of the joint reconstruction 29
5.8 Example: five targets on a mountain 29

6.1 The DEM in radar geometry 31
6.2 The model of the 2-D joint reconstruction 31
6.3 Examples of the 2-D joint reconstruction 33

7.1 Median highest ambiguity of every method 34

7.2 All the methods on a realistic scene 36
7.3 SATA against the joint reconstruction, every run 37
7.4 1-D sweep by number of channels 37
7.5 Every run of the sweeps 38
7.6 Joint reconstruction with 2, 3, 4 and 8 channels 39
7.7 Example of the 1-D sweep 41
7.8 Sensitivity to DEM errors 41
7.9 SRTM DEM of the Sierra Nevada and the five named lines 42
7.10 Real terrain: Yosemite Valley 43
7.11 Real terrain: eastern Sierra escarpment 44
7.12 Real terrain: High Sierra 45
7.13 Real terrain: zoom on one target 45
7.14 Real terrain: 20 targets per line 46
7.15 Real terrain: 200 random lines 47

List of Tables

2.1 Parameters of the simulated system. 3

5.1 Objects of the joint model 25
5.2 How the topography enters the joint model 27

7.1 Comparison of all methods 34
7.2 2-D sweep by number of channels 38
7.3 Real terrain: the five named lines 40

Symbols and acronyms

Symbol	Meaning
λ	wavelength (0.25 m)
v	platform velocity
r_0	slant range of closest approach of the reference point
θ	look angle ($\approx 20°$)
N_{rx}	number of receive channels (one transmitter, N_{rx} receivers)
$b_{\mathrm{at},i}$, $b_{\mathrm{xt},i}$	along-track and cross-track baseline of receiver i relative to the transmitter
PRF	full (reconstructed) pulse repetition frequency, 2 000 Hz
$\mathrm{PRF}_{\mathrm{op}}$	PRF of each channel, $\mathrm{PRF}/N_{\mathrm{rx}}$
B_{a}	processed azimuth (Doppler) bandwidth, 1 235 Hz
$C_0, C_1, C_2, \Delta t$	coefficients of the bistatic minus monostatic path of a channel (Section 2.3)
$\Delta C_{0,i}(x)$	topographic residual of channel i: C_0 of a scatterer at height $h(x)$ minus C_0 of the reference point
X	ground distance between parcels folded onto the same Doppler bin, $X \approx r_0 \lambda \mathrm{PRF}_{\mathrm{op}}/(2v)$
δx	resolution of one SATA sub-aperture, $\sqrt{\lambda r_0/2}$
$H_i(f)$	transfer function of channel i relative to the monostatic reference
$g(x)$	focused full-rate azimuth line (unknown of the joint reconstruction)
μ	Tikhonov regularisation parameter
Acronym	
CGLS	conjugate gradients for least squares
DEM	digital elevation model
DPCA	displaced phase centre antenna (uniform effective sampling)
HRWS	high-resolution wide-swath
IRF	impulse response function
SATA	sub-aperture topography- and aperture-dependent algorithm

1 Introduction

1.1 Context

A conventional SAR cannot improve its azimuth resolution and its swath width at the same time: a fine azimuth resolution needs a short antenna and therefore a wide Doppler bandwidth, which must be sampled with a high PRF, while a wide swath needs a long interval between pulses. Multichannel systems break this trade-off by receiving each pulse with several apertures displaced along track [1, 2]. Each channel is sampled at a PRF below the Doppler bandwidth and is therefore aliased, but the set of channels samples the azimuth signal at the full rate, and the unambiguous spectrum is recovered by an *azimuth reconstruction*: a bank of filters that combines the channels so that the aliased parts cancel.

Multistatic formations distribute the receivers over several platforms. They lift the limit set by the length of a single antenna and offer flexibility, but the baselines between platforms are no longer small [3, 4]. In particular, a *cross-track* baseline makes the bistatic range history of a scatterer depend on its height. The reconstruction filters are derived for a reference geometry (here: the reference point of the scene at height $h_0 = 0$); a scatterer at another height carries a residual phase that is different in each channel. The filters then no longer match the data, the aliased parts no longer cancel and azimuth ambiguities appear in the image.

The same kind of height-dependent mismatch is known from airborne motion compensation, where the compensation is computed for a reference height and leaves residual errors for the rest of the topography. Prats et al. proposed the sub-aperture topography- and aperture-dependent (SATA) algorithm to correct it with an external DEM [5, 6]: the azimuth signal is split into short sub-apertures, each Doppler frequency of a sub-aperture is associated with a ground position, and the phase residual of that position, read from the DEM, is removed.

1.2 Objective

The fellowship topic is to investigate whether SATA can correct the azimuth ambiguities caused by topography and cross-track baselines in multistatic SAR systems. During the work the question was made precise as:

1. How large are the ambiguities caused by unaccounted topography, as a function of the terrain, the baselines and the number of channels?
2. How much of them can be removed using *only the system parameters and a DEM*, with no knowledge of where the scatterers are?

The second condition is strict on purpose. Several intermediate methods tested during the project reached the reference level only because their correction maps were built from the positions of the simulated targets; such methods are not applicable to real data and are not counted as solutions here (Section 4.2.6).

1.3 Contributions

- A numerical simulator of the multistatic acquisition and of the multichannel reconstruction in one dimension (azimuth) and in two dimensions (azimuth and range, 3-D scenes), with along- and cross-track baselines and arbitrary terrain (Chapters 2 and 3).
- An analysis of SATA applied to each aliased channel before the reconstruction, showing when it is exact and predicting the residual it leaves on general terrain (Chapter 4).
- A joint reconstruction that places the DEM phase inside the inversion, first per sub-aperture and finally over the whole line, solved by regularised least squares with conjugate gradients (Chapter 5), and its 2-D extension (Chapter 6).
- Systematic sweeps over terrains, numbers of channels and baselines, and an audit of every method with a DEM only (Chapter 7).

1.4 Outline

Chapter 2 introduces the acquisition geometry and the topographic residual ΔC_0. Chapter 3 reviews the multichannel reconstruction and explains how topography produces ambiguities. Chapter 4 describes SATA and its limit for aliased channels. Chapter 5 presents the joint reconstruction and Chapter 6 its 2-D version. Chapter 7 collects the results and Chapter 8 concludes. The simulator, the methods and the scripts that produce every figure of this report are available at github.com/lucas-ayub/dlr-project.

2 Acquisition geometry and signal model

2.1 Formation and parameters

One platform transmits and receives (channel 0); the other $N_{\mathrm{rx}} - 1$ receivers fly behind it. Receiver i is displaced by $b_{\mathrm{at},i}$ along track and by $b_{\mathrm{xt},i}$ across track (horizontally, towards the scene) relative to the transmitter (Figure 2.1). The along-track baselines satisfy the displaced phase centre (DPCA) condition (Figure 2.2), so that the effective phase centres of the N_{rx} channels are uniformly spaced and the multichannel sampling is uniform at the full PRF; the study is about the cross-track baselines, which are drawn at random in $[0, {b_{\mathrm{xt}}}^{\max}]$ (with the largest equal to ${b_{\mathrm{xt}}}^{\max}$) or taken from fixed sets. Channel 0 is always the transmitter ($b_{\mathrm{at}} = b_{\mathrm{xt}} = 0$). The parameters are listed in Table 2.1.

Table 2.1: Parameters of the simulated system.

Wavelength λ	0.25 m
Platform altitude H	720 km
Slant range of the reference point r_0	766.2 km
Look angle θ	20°
Platform velocity / effective velocity	7 688.5 m/s / 7 408.5 m/s
Antenna length	12 m
Processed Doppler bandwidth B_{a}	1 235 Hz
Full PRF	2 000 Hz
Channel PRF $\mathrm{PRF}_{\mathrm{op}} = \mathrm{PRF}/N_{\mathrm{rx}}$	1 000 Hz to 250 Hz ($N_{\mathrm{rx}} = 2$ to 8)
DPCA along-track spacing $2v/\mathrm{PRF}$	7.69 m
Cross-track baselines ${b_{\mathrm{xt}}}^{\max}$	100 m, 300 m and 500 m, and fixed sets

2.2 Signal model

The simulator works with point targets. In the 1-D study the targets lie on a terrain profile $h(x)$ along azimuth; each target is placed on the iso-range surface of the reference point (same monostatic slant range r_0), so that only its height changes the bistatic geometry. The reference point is at $x = 0$, $h = 0$. In the 2-D study (Chapter 6) the terrain is a surface $h(x, y)$ and the data are range compressed. The azimuth signal of channel i is

$$s_i(t) = \sum_p a_p \, w(t - t_p) \, \exp\left(-\mathrm{j} \frac{2\pi}{\lambda} \left[r_{\mathrm{T},p}(t) + r_{\mathrm{R},i,p}(t) \right] \right), \tag{2.1}$$

with a_p the complex amplitude of scatterer p, w the antenna pattern and $r_{\mathrm{T}} + r_{\mathrm{R},i}$ the bistatic path (Figure 2.1).

2.3 Channel transfer functions and the coefficients C_0, C_1, C_2, Δt

The reconstruction (Chapter 3) needs, for every channel i, its transfer function $H_i(f)$: the factor by which the Doppler spectrum of channel i differs from the spectrum the monostatic system (transmitter and receiver at the transmitter) would record for the same scatterer. This

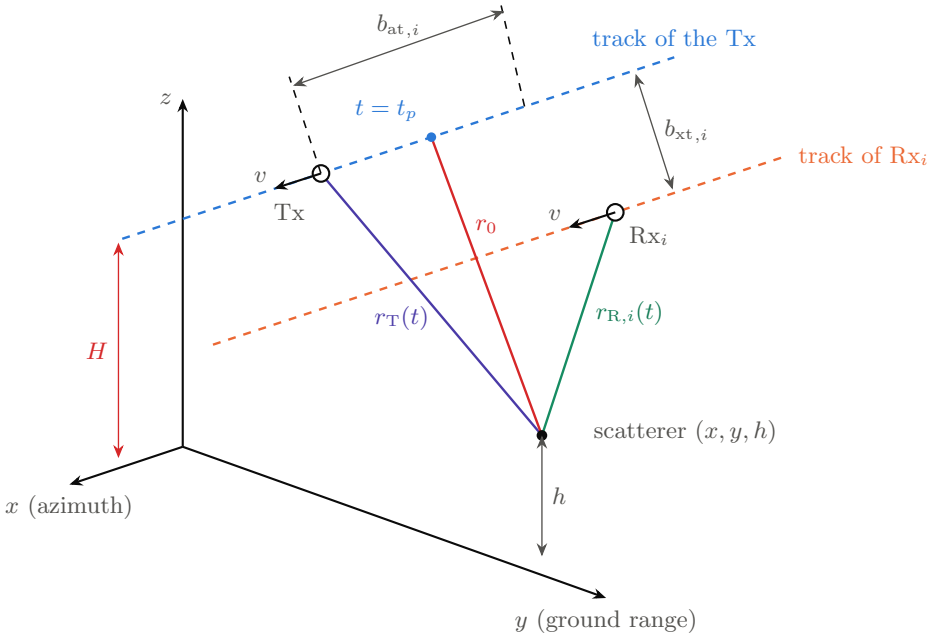


Figure 2.1: Geometry of one channel. The transmitter (Tx) and receiver i (Rx$_i$) fly on parallel tracks at the same altitude H, separated by $b_{\mathrm{at},i}$ along track and $b_{\mathrm{xt},i}$ across track. The echo of channel i travels $r_{\mathrm{T}}(t) + r_{\mathrm{R},i}(t)$. The time of closest approach t_p and the range r_0 fix the position of the scatterer on the line; its height h is what the topography adds. Not to scale.

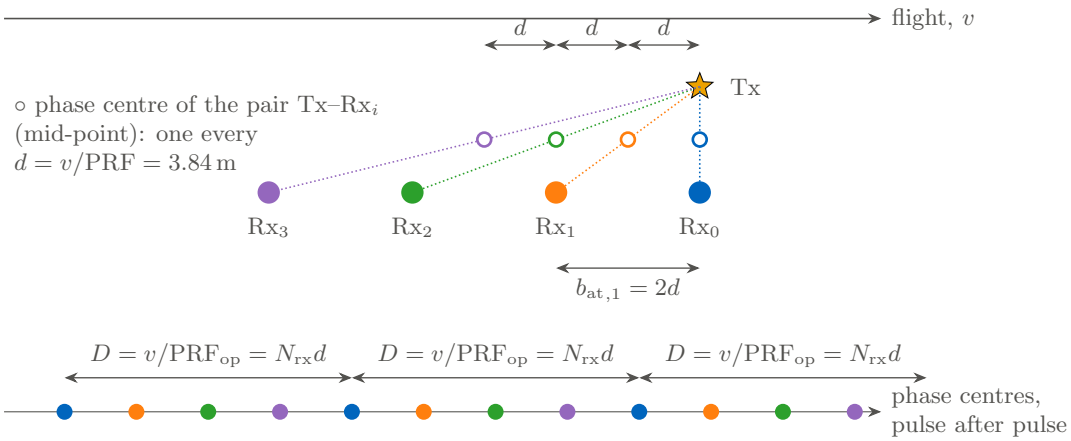


Figure 2.2: DPCA arrangement seen from above ($N_{\mathrm{rx}} = 4$). Receiver i trails the transmitter by $b_{\mathrm{at},i} = 2id$, so the phase centre of the pair (the mid-point) trails it by id. At every pulse the four phase centres occupy four equally spaced positions; between pulses the platform moves $D = N_{\mathrm{rx}}d$ and the pattern repeats. Together the channels sample the ground every d, like a single channel at the full PRF.

section derives H_i from the range histories of the reference point; the derivation follows the Doppler-domain approach of [3, 4].

Step 1: two parabolas (time domain). Let $M(t) = 2r_{\mathrm{T}}(t)$ be the monostatic path and $B_i(t) = r_{\mathrm{T}}(t) + r_{\mathrm{R},i}(t)$ the bistatic path of channel i, both for the reference point. Over the aperture each is well described by a parabola around its own minimum:

$$M(t) \approx M_{\min} + m_2\,(t - t_M)^2, \qquad B_i(t) \approx B_{i,\min} + b_{2,i}\,(t - t_{B,i})^2, \tag{2.2}$$

where $M_{\min}$, $B_{i,\min}$ are the smallest path lengths [m], t_M, $t_{B,i}$ the times at which they occur [s] and m_2, $b_{2,i}$ the curvatures [m/s^2]. The corresponding azimuth signals are $s_M(t) = w(t)\exp[-\mathrm{j}2\pi M(t)/\lambda]$ and $s_i(t) = w(t)\exp[-\mathrm{j}2\pi B_i(t)/\lambda]$.

Step 2: from time to Doppler (stationary phase). For a signal $\exp[-\mathrm{j}2\pi P(t)/\lambda]$ with a parabolic path $P(t) = P_{\min} + p_2(t - t_P)^2$, the instantaneous frequency is

$$f(t) = -\frac{1}{\lambda}\frac{\mathrm{d}P}{\mathrm{d}t} = -\frac{2p_2}{\lambda}(t - t_P) \quad \Longrightarrow \quad t_f = t_P - \frac{\lambda f}{2p_2}, \tag{2.3}$$

so every Doppler frequency f is produced at one time t_f (the same relation as the ruler (3.3) of Chapter 3). By the principle of stationary phase the phase of the spectrum at f is the phase of the signal at t_f minus $2\pi f t_f$. With $P(t_f) = P_{\min} + \lambda^2 f^2/(4p_2)$,

$$\Psi_P(f) = -\frac{2\pi}{\lambda}P(t_f) - 2\pi f t_f = -2\pi\Big[\frac{P_{\min}}{\lambda} + t_P\,f - \frac{\lambda f^2}{4p_2}\Big]. \tag{2.4}$$

The amplitude of the spectrum is set by the antenna pattern and is the same for the two paths.

Step 3: the ratio of the two spectra (Doppler domain). Applying (2.4) to B_i and to M and taking the difference of the phases gives the transfer function of channel i:

$$H_i(f) = \mathrm{e}^{\mathrm{j}[\Psi_{B_i}(f) - \Psi_M(f)]} = \exp\Big(-\mathrm{j}2\pi\Big[\frac{C_{0,i}}{\lambda} + (\Delta t_i - C_{1,i})\,f + C_{2,i}\,\lambda f^2\Big]\Big), \tag{2.5}$$

with

$$C_{0,i} = B_{i,\min} - M_{\min}, \qquad \Delta t_i - C_{1,i} = t_{B,i} - t_M, \qquad C_{2,i} = \frac{1}{4m_2} - \frac{1}{4b_{2,i}}. \tag{2.6}$$

In the simulator the two paths are sampled at the full PRF: Δt_i is the time between the samples where M and B_i are smallest and $C_{1,i}$ the sub-sample correction given by the linear terms of the two fitted parabolas, so that $\Delta t_i - C_{1,i}$ is the time between the two minima.

Physical meaning. Each coefficient corresponds to one effect of the bistatic geometry:

- $C_{0,i}$ [m], the difference of the shortest paths, is a *constant phase* $2\pi C_{0,i}/\lambda$. It is set by the cross-track baseline, and it is the coefficient that changes with the height of the scatterer (Section 2.4).
- $\Delta t_i - C_{1,i}$ [s], the delay between the two minima, is a *linear phase* in Doppler, i.e. a shift in time. It is set by the along-track baseline ($\approx b_{\mathrm{at},i}/(2v)$) and is what makes the channels sample the azimuth signal at displaced positions (DPCA).
- $C_{2,i}$ [s^2/m], the difference of the inverse curvatures, is a *quadratic phase*: a small residual azimuth chirp.

For a scatterer other than the reference point the same derivation holds with its own paths; what changes, to first order, is $C_{0,i}$ (Section 2.4).

Figure 2.3 shows the four coefficients of the reference point. C_0 is the difference of the closest-approach distances of the receiver and of the transmitter, $\approx -b_{\mathrm{xt}} \sin\theta$ ($-171\,\mathrm{m}$ for $b_{\mathrm{xt}} = 500\,\mathrm{m}$), and does not depend on b_{at}; $\Delta t = b_{\mathrm{at}}/(2v)$ (0.5 ms, 1.0 ms and 1.5 ms for the three DPCA receivers) and does not depend on b_{xt}; C_1 is below $0.2\,\mu\mathrm{s}$ and C_2 below $4 \times 10^{-7}\,\mathrm{s}^2/\mathrm{m}$, which at the edge of the processed band, $f = B_{\mathrm{a}}/2$, give phases of $2\pi C_1 f < 1 \times 10^{-3}$ rad and $2\pi C_2 \lambda f^2 \approx 0.2$ rad. Figure 2.4 checks the parabolic model (2.2) for one target: the bistatic minus the monostatic path varies by a few millimetres over the aperture (5.6 mm for $b_{\mathrm{xt}} = 300\,\mathrm{m}$) and the parabola reproduces it within 10^{-4} mm.

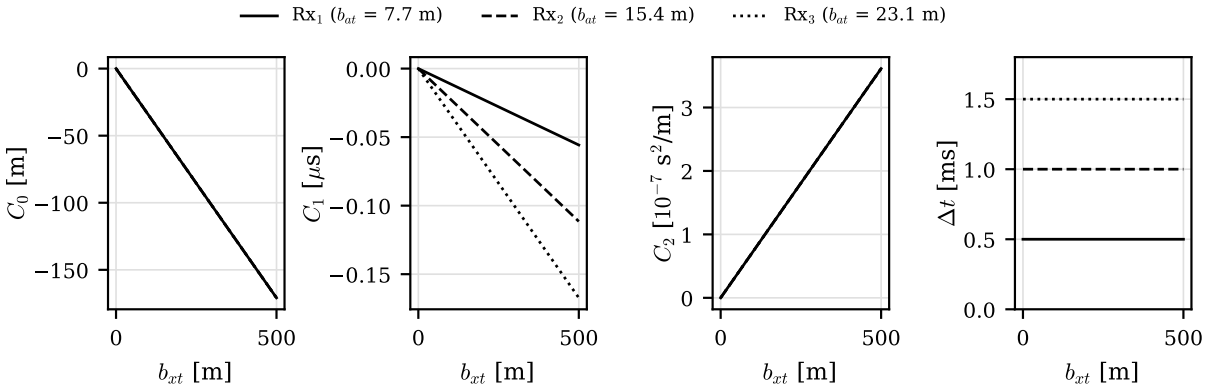


Figure 2.3: C_0, C_1, C_2 and Δt of the reference point for the receivers 1, 2 and 3 of the DPCA formation ($b_{\mathrm{at}} = 7.7\,\mathrm{m}$, 15.4 m and 23.1 m), as a function of the cross-track baseline b_{xt} (the same for the three receivers). The three curves of C_0 and of C_2 coincide.

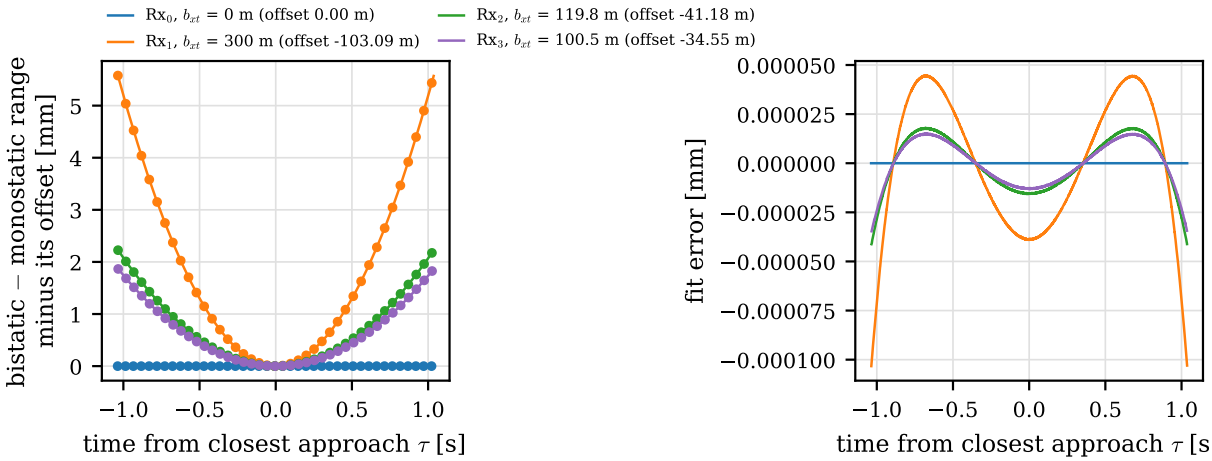


Figure 2.4: One target at $x = 0$, $h = 500\,\mathrm{m}$, four channels with $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m. Left: bistatic minus monostatic path of every channel against the time from closest approach, minus its constant offset (C_0, given in the legend); dots: simulated geometry; lines: the fitted parabola. Right: difference between the simulated curve and the parabola.

2.4 The topographic residual ΔC_0

For a scatterer at height h the coefficient C_0 of channel i differs from the one of the reference point. The difference is the *topographic residual*

$$\Delta C_{0,i}(h) = C_{0,i}(\text{scatterer}) - C_{0,i}(\text{reference point}) \approx -\frac{b_{\text{xt},i}\, h}{r_0 \tan\theta}, \qquad \varphi_i(h) = \frac{2\pi}{\lambda}\Delta C_{0,i}(h), \quad (2.7)$$

i.e. the interferometric phase of the cross-track baseline; it vanishes for the transmitter channel ($b_{\text{xt}} = 0$), since the targets are on the iso-range surface. The simulator computes $\Delta C_{0,i}$ exactly from the range histories; the approximation (2.7) agrees within about 2 %. For a scatterer at 500 m and $b_{\text{xt}} = [63.7, 27.0, 4.1, 1.7]$ m the residuals are $[-113.9, -48.2, -7.3, -3.0]$ mm, i.e. up to half a wavelength. The other coefficients change much less with height: the topography acts as a channel-dependent *constant* phase over the aperture of each scatterer, a fact used by the joint reconstruction (Chapter 5). Figure 2.5 quantifies this: up to 3 km of height and $b_{\text{xt}} = 500$ m, ΔC_0 follows (2.7) within a few per cent, and at the edge of the processed band the phase produced by ΔC_0 (about 80 rad at 3 km with $b_{\text{xt}} = 300$ m) is four orders of magnitude larger than the one produced by ΔC_2 (4×10^{-3} rad) and seven larger than the one produced by ΔC_1 (4×10^{-6} rad); Δt does not change with height.

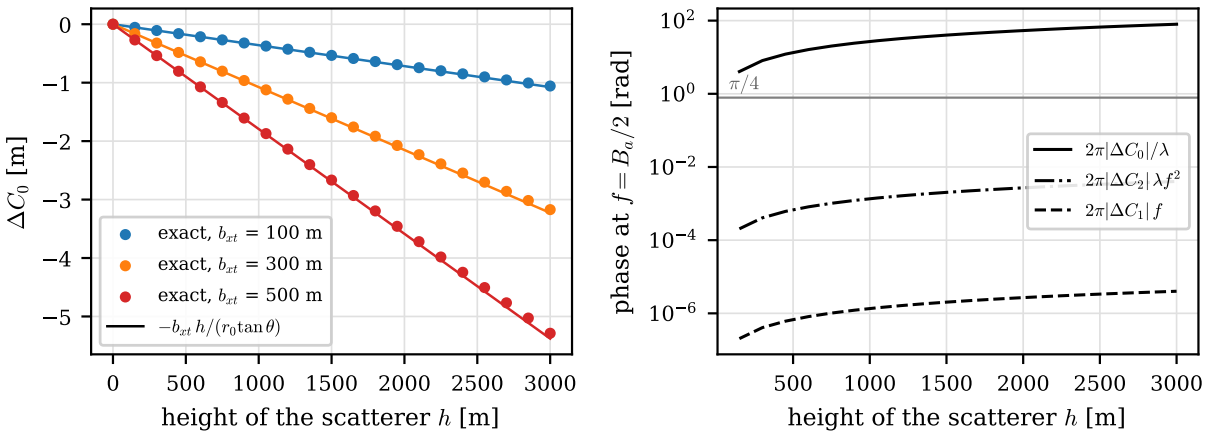


Figure 2.5: Change of the coefficients of receiver 1 when the scatterer is at height h instead of the reference point. Left: ΔC_0 (dots: exact, from the range histories; lines: (2.7)) for three cross-track baselines. Right ($b_{\text{xt}} = 300$ m): phase produced by each term of (2.5) at the edge of the processed band, $f = B_{\text{a}}/2$, on a logarithmic scale.

Figure 2.6 shows where (2.7) comes from. Moving from the reference point to a scatterer at height h on the same range arc of the transmitter is a displacement $s = h/\sin\theta$ perpendicular to the line of sight of the transmitter, which does not change its path. The line of sight of receiver i is rotated by $\Delta\theta = b_{\text{xt},i}\cos\theta/r_0$, so its path changes by $s\,\Delta\theta = b_{\text{xt},i}h/(r_0\tan\theta)$.

The sensitivity is $\partial\varphi_i/\partial h = 2\pi b_{\text{xt},i}/(\lambda r_0 \tan\theta) \approx 9.0 \times 10^{-5}\, b_{\text{xt},i}$ rad/m, which also sets the tolerance to DEM errors (Section 7.5).

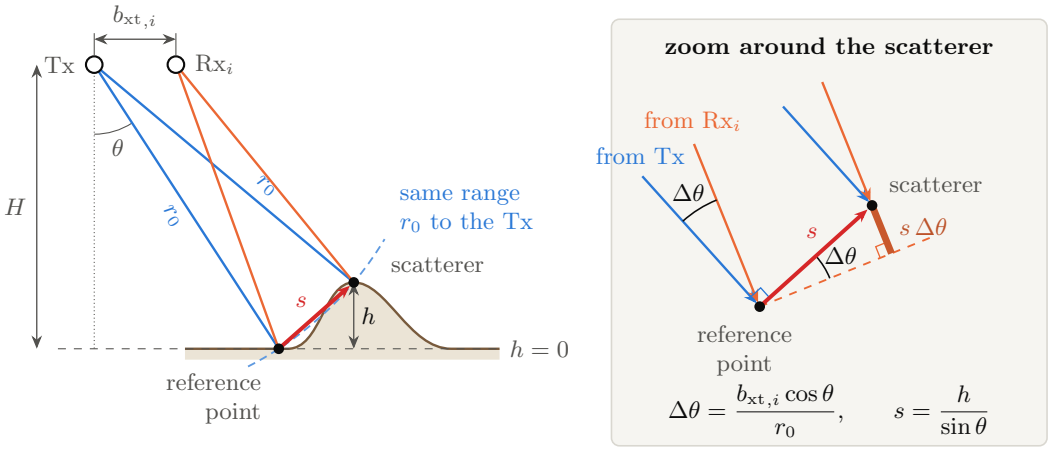


Figure 2.6: Left: plane perpendicular to the track (not to scale). The reference point and the scatterer are at the same range r_0 from the transmitter, at different heights: climbing h along the arc is a displacement $s = h/\sin\theta$. Right: zoom around the scatterer. Since $s \ll r_0$ the lines of sight of each platform to the two points are parallel; s is perpendicular to the line of sight of the transmitter, but the one of the receiver is rotated by $\Delta\theta$, and its path becomes $s\,\Delta\theta$ shorter (thick orange segment).

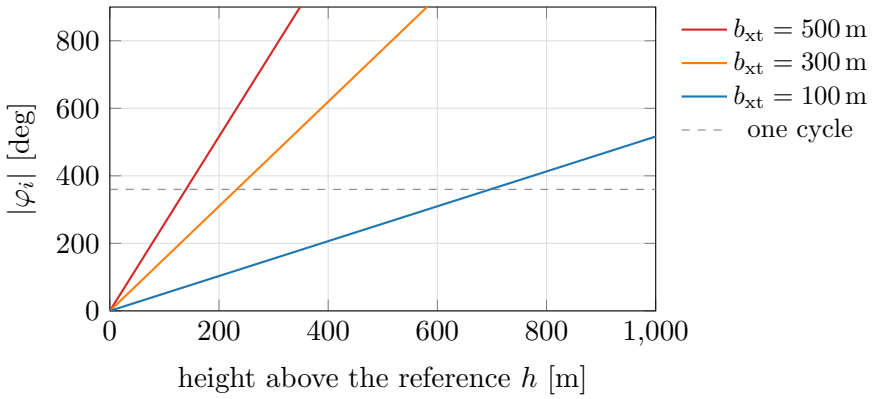


Figure 2.7: Phase of the topographic residual (2.7) against height. With $b_{\mathrm{xt}} = 300\,\mathrm{m}$ the sensitivity is 1.55° per metre of height: a 230 m hill already produces a full cycle.

3 Multichannel azimuth reconstruction and topography

3.1 Aliased channels and folded parcels

Each channel is sampled at $\mathrm{PRF}_{\mathrm{op}} = \mathrm{PRF}/N_{\mathrm{rx}}$, below the Doppler bandwidth $B_{\mathrm{a}} = 1\,235\,\mathrm{Hz}$. In the Doppler domain of the channel, every frequency $\phi \in [-\mathrm{PRF}_{\mathrm{op}}/2, \mathrm{PRF}_{\mathrm{op}}/2)$ is the sum of N_{rx} *parcels* of the full-rate spectrum U, taken at

$$f_q = \phi + q\,\mathrm{PRF}_{\mathrm{op}} - \tfrac{N_{\mathrm{rx}}-1}{2}\mathrm{PRF}_{\mathrm{op}}, \qquad q = 0, \dots, N_{\mathrm{rx}} - 1, \tag{3.1}$$

each weighted by the transfer function H_i of the channel relative to the monostatic reference:

$$S_i(\phi) = \sum_q H_i(f_q)\, U(f_q) \quad \Longleftrightarrow \quad \boldsymbol{s}(\phi) = \mathbf{H}(\phi)\, \boldsymbol{u}(\phi), \tag{3.2}$$

with $\boldsymbol{s}$ the N_{rx} channel spectra and $\boldsymbol{u}$ the N_{rx} parcels. The standard reconstruction [1, 2] inverts this square system at every ϕ, $\hat{\boldsymbol{u}} = \mathbf{P}(\phi)\boldsymbol{s}(\phi)$ with $\mathbf{P} = \mathbf{H}^{-1}$, and places the parcels at their true Doppler frequency to obtain the full-rate signal. The entries of $\mathbf{H}$ are the transfer functions (2.5) of the reference point, $H_{iq}(\phi) = H_i(f_q)$.

Through the Doppler–angle relation, a Doppler frequency f seen from the platform at azimuth position x_c corresponds to the ground position

$$x(f) = x_c + r_0 \tan\!\big(\arcsin(\lambda f/2v)\big) \approx x_c + \frac{\lambda r_0}{2v}\, f, \tag{3.3}$$

which we call the *ruler*. Parcels one $\mathrm{PRF}_{\mathrm{op}}$ apart therefore come from ground positions

$$X \approx \frac{\lambda r_0\, \mathrm{PRF}_{\mathrm{op}}}{2v} \qquad (6.23\,\mathrm{km} \text{ for } \mathrm{PRF}_{\mathrm{op}} = 500\,\mathrm{Hz}), \tag{3.4}$$

apart, i.e. 0.81 s of azimuth time. This distance is the key length of the problem: one bin of an aliased channel mixes the echoes of the terrain at x, $x \pm X$, $x \pm 2X$, … (Figure 3.1).

Figure 3.2 shows the two sides of this for one target on flat terrain and four channels. A single channel at $\mathrm{PRF}_{\mathrm{op}} = 500\,\mathrm{Hz}$ sees the processed band $B_{\mathrm{a}} = 1\,235\,\mathrm{Hz}$ folded two to three times onto its own band; the standard reconstruction combines the four channels and recovers the full-rate spectrum of the reference.

3.2 How topography creates ambiguities

For a scatterer at height h, channel i carries the extra constant phase $\mathrm{e}^{-\mathrm{j}\varphi_i}$ of (2.7), so the data are $\boldsymbol{s} = \mathbf{D}\mathbf{H}\boldsymbol{u}$ with $\mathbf{D} = \mathrm{diag}(\mathrm{e}^{-\mathrm{j}\varphi_i})$, and the standard reconstruction returns

$$\hat{\boldsymbol{u}} = \mathbf{H}^{-1}\mathbf{D}\,\mathbf{H}\,\boldsymbol{u} = \mathbf{L}\,\boldsymbol{u}. \tag{3.5}$$

In the DPCA configuration $\mathbf{H}$ is, up to diagonal factors, a discrete Fourier matrix over the channels, and the leakage matrix becomes Toeplitz,

$$L_{q'q} = c_{q'-q}, \qquad c_m = \frac{1}{N_{\mathrm{rx}}} \sum_{i=0}^{N_{\mathrm{rx}}-1} \exp\!\Big[\mathrm{j}\Big(\frac{2\pi i m}{N_{\mathrm{rx}}} - \varphi_i\Big)\Big]. \tag{3.6}$$

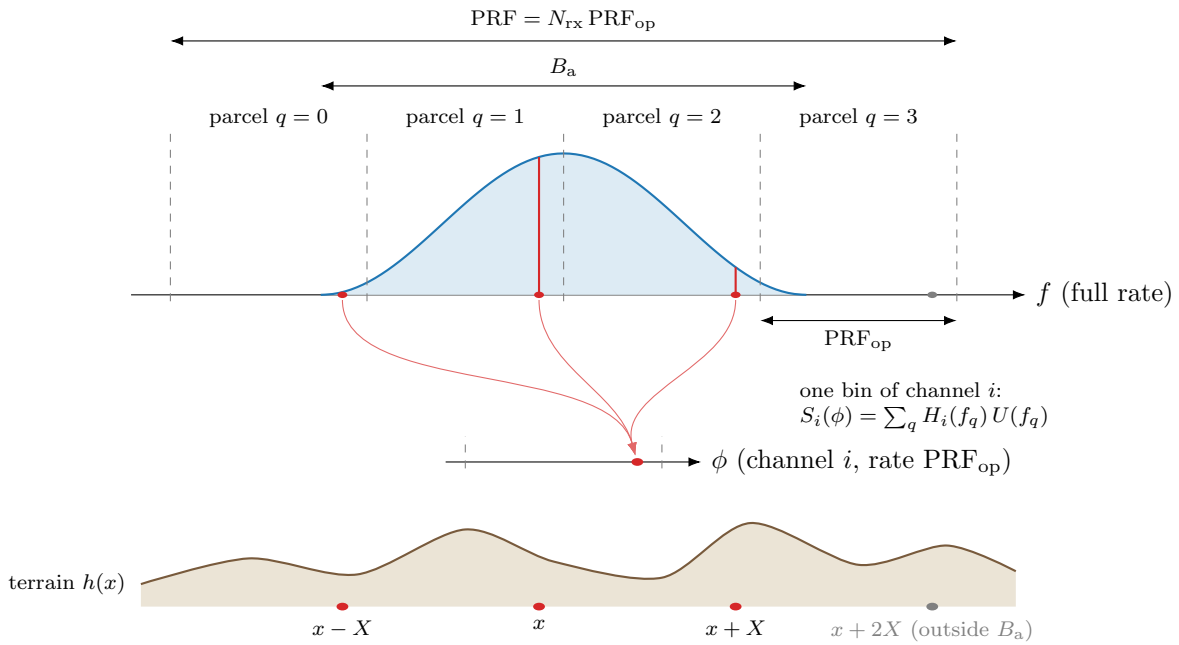


Figure 3.1: Folding of the azimuth spectrum in an aliased channel ($N_{rx} = 4$). The four parcels of one channel bin come, through the ruler (3.3), from ground positions X apart, which in general have different heights.

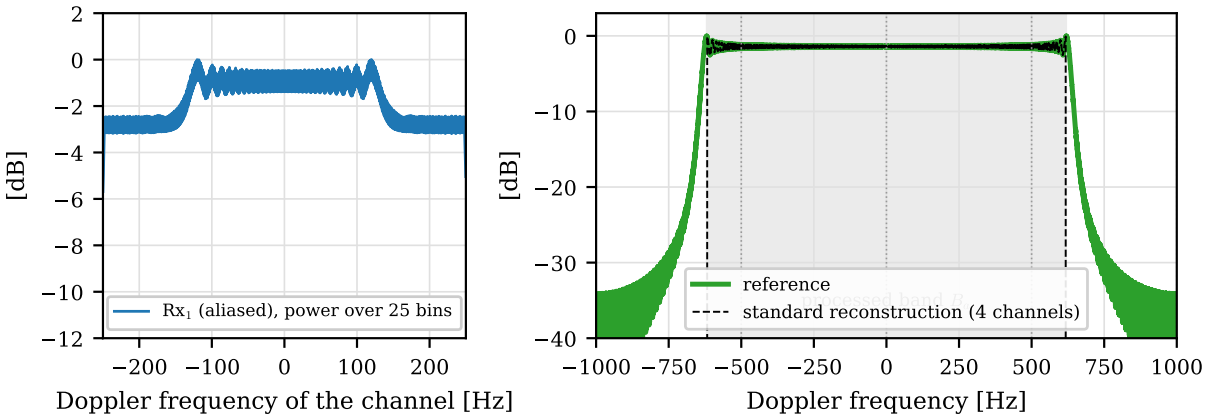


Figure 3.2: One target at $h = 0$, four channels ($b_{xt} = [0, 300, 119.8, 100.5]$ m). Left: power spectrum of one channel at PRF$_{op}$ = 500 Hz (averaged over 25 bins): the folded parcels overlap. Right: spectrum of the reference and of the standard reconstruction at the full PRF; the grey band is the processed band B_a and the dotted lines are the borders of the four parcels, PRF$_{op}$ apart.

c_0 is the gain of the true parcel and c_m, $m \neq 0$, the amplitude that parcel q leaks into parcel $q + m$, which appears in the image as an ambiguity displaced by mX ($0.81\,\mathrm{s}{\cdot}m$). If all channels have the same phase, $c_m = 0$ for $m \neq 0$ and the reconstruction is exact; with cross-track baselines the phases differ and the ambiguity level grows with the spread of $b_{\mathrm{xt},i}h$. Weighted by the fraction of the target bandwidth that falls in the leaking parcels, $|c_m|$ predicts the ambiguity of the uncorrected reconstruction within 0.15 dB for heights above 50 m.

Figure 3.3 shows the effect on one target. At $h = 0$ the channels carry no residual and the standard reconstruction is the reference (up to the band limitation to B_a); at 500 m copies of the target appear at multiples of the ambiguity time, 0.81 s. Figure 3.4 gives the highest ambiguity against the height of the target for three sets of baselines: it grows by 20 dB per decade of height (for small phases $|c_m|$ is proportional to φ_i, hence to $b_{\mathrm{xt}}h$) and saturates between -15 dB and -5 dB once the phases of the channels are of the order of a radian, i.e. above about 100 m with ${b_{\mathrm{xt}}}^{\max} = 300$ m.

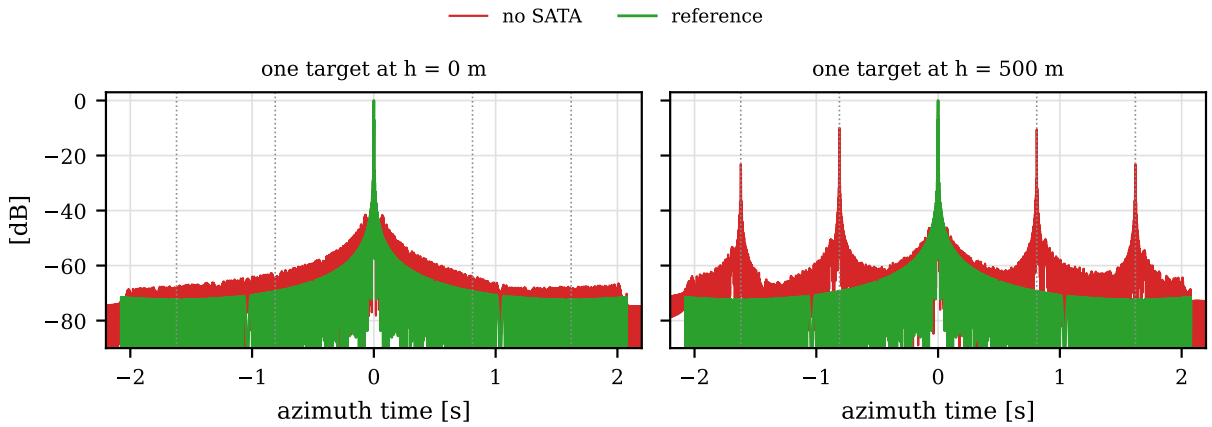


Figure 3.3: Azimuth impulse response of one target ($b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m) after the standard reconstruction (no SATA) and of the reference. Dotted lines: the ambiguity positions, ± 1 and ± 2 times 0.81 s. Left: target at the reference height, the two responses coincide. Right: target at 500 m.

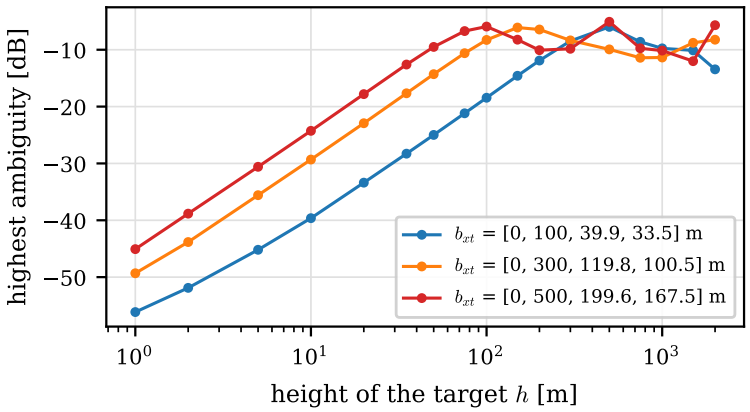


Figure 3.4: Highest ambiguity of the standard reconstruction (no SATA) for one target at height h, four channels, three sets of cross-track baselines. At $h = 0$ the three sets give -63.6 dB, the level of the reference limited to the processed band B_a.

Two consequences recur in the rest of the report:

- The amplitudes c_m are the discrete Fourier transform of the channel phases. Baselines for

which several φ_i coincide modulo 2π, or that are equally spaced, make the system badly conditioned at certain heights: the channels no longer separate the parcels.

- With two channels there is no redundancy at all: two parcels, two equations.

3.3 Ambiguity metric

All results are reported as the *highest ambiguity*: the largest value of the azimuth impulse response, in dB relative to the peak of the reference, inside ± 20 ms of the ambiguity positions of every target (± 1 and ± 2 times the ambiguity distance of the channel PRF). The reference is the monostatic full-rate signal of the same scene, processed with the same band limitation; its own value (sidelobes of neighbouring targets, band edge) is the floor that any reconstruction can reach. Repeating the measurement with orders up to ± 5 for six and eight channels did not change the conclusions.

4 SATA before the reconstruction

4.1 The SATA algorithm

SATA was developed for airborne motion compensation [5, 6]. Its idea is that a short piece of the azimuth signal is a low-resolution image of the ground: inside a sub-aperture, each Doppler frequency belongs to one ground position through the ruler (3.3). The topography-dependent phase can therefore be removed frequency by frequency, with the value of the DEM at the position that frequency points to.

4.1.1 Why a sub-aperture is an image

Take a sub-aperture of duration T centred on the platform position x_c. A scatterer at x_p has, at that moment, the Doppler $f \approx 2v(x_p - x_c)/(\lambda r_0)$, and during the sub-aperture its Doppler sweeps $B_T = 2v^2T/(\lambda r_0)$. In the spectrum of the sub-aperture the scatterer is therefore a peak at the frequency that the ruler maps back to x_p. Two effects blur this image of the ground:

- the sweep B_T, which through the ruler covers a ground length $\frac{\lambda r_0}{2v} B_T = vT$, the length of the sub-aperture itself;
- the frequency resolution $1/T$, which covers $\frac{\lambda r_0}{2v} \cdot \frac{1}{T} = \frac{\lambda r_0}{2vT}$ of ground.

The first grows and the second shrinks with T; they are equal for

$$vT = \frac{\lambda r_0}{2vT} \quad \Longrightarrow \quad \delta x = vT = \sqrt{\frac{\lambda r_0}{2}} = 309.5\,\mathrm{m}, \tag{4.1}$$

the ground resolution of SATA. At $\mathrm{PRF}_{\mathrm{op}} = 500\,\mathrm{Hz}$ this is a sub-aperture of 20 samples (40 ms). Since the residual $\Delta C_{0,i}$ of a scatterer is a constant phase over its whole aperture (Chapter 2), multiplying a bin by the phase of the position it points to removes the topographic phase of whatever lies in that cell of about 310 m.

4.1.2 The steps

For one channel i (Figure 4.1):

1. cut the azimuth line into sub-apertures b of $T = 20$ samples, with a triangular window w and a hop of 10 samples (50 % overlap, so that the windows add up to a constant);
2. Fourier transform each sub-aperture with zero padding to $N_{zp} = 128$ points: bins f_m in the band of the channel;
3. give each bin a ground position with the ruler, $x_{b,m} = x_{c,b} + r_0 \tan\big(\arcsin(\lambda f_m/2v)\big)$, read the residual of that position from the DEM and rotate the bin,

$$S_b(f_m) \;\leftarrow\; S_b(f_m)\, \exp\Big(+\mathrm{j}\frac{2\pi}{\lambda}\, \Delta C_{0,i}\big(x_{b,m}\big)\Big); \tag{4.2}$$

4. inverse transform, keep the 20 samples, overlap-add the sub-apertures and divide by the sum of the windows.

The map $\Delta C_{0,i}(x)$ is computed once per channel from the DEM: the height $h(x)$ is converted into the residual of (2.7) with the bistatic geometry of channel i (through a table of $\Delta C_{0,i}$ against height). With a zero map the kernel is the identity ($< 10^{-15}$), and with a constant map it rotates the line by that phase.

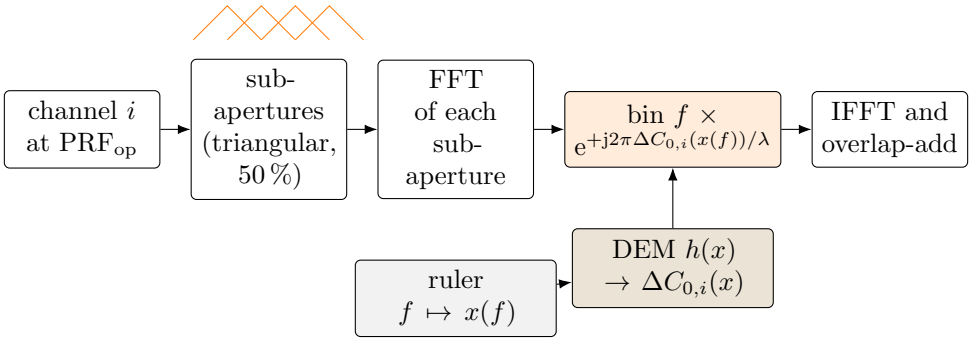


Figure 4.1: The SATA kernel applied to one channel. The only external information is the DEM; the ruler associates each sub-aperture bin with a ground position.

4.2 SATA on each channel, then the standard reconstruction

The natural way to use SATA in a multichannel system is to correct every channel with its own $\Delta C_{0,i}$ and then run the standard reconstruction. The difficulty is step 3: the kernel needs the Doppler frequency of each bin, and in an aliased channel a bin $\phi \in [-\mathrm{PRF}_{\mathrm{op}}/2, \mathrm{PRF}_{\mathrm{op}}/2)$ holds the N_{rx} parcels $f_q = \phi + q\mathrm{PRF}_{\mathrm{op}} - \frac{N_{\mathrm{rx}}-1}{2}\mathrm{PRF}_{\mathrm{op}}$ (Figure 3.1). Some frequency must be chosen to label the bin, and the two variants differ only in that choice (Figure 4.2).

4.2.1 Whole band

Each bin is labelled with the Doppler it has in the channel, $f = \phi$, as if the channel were not aliased. Each channel is corrected once with (4.2), and the corrected channels go through the standard reconstruction (3.2). The cost is one SATA per channel. The correction is right only for the parcel that actually lies at $f = \phi$, i.e. for the part of the target spectrum inside $[-\mathrm{PRF}_{\mathrm{op}}/2, \mathrm{PRF}_{\mathrm{op}}/2)$; the other parcels of the bin receive the phase of a position X, $2X, \ldots$ away.

4.2.2 Per sub-band

The output spectrum is divided into N_{rx} sub-bands of width $\mathrm{PRF}_{\mathrm{op}}$, centred at $f_k = \left(k - \frac{N_{\mathrm{rx}}}{2} + \frac{1}{2}\right)\mathrm{PRF}_{\mathrm{op}}$, $k = 0, \ldots, N_{\mathrm{rx}} - 1$. For every k:

1. each bin is labelled with its alias inside sub-band k,

$$f_k(\phi) = f_k - \frac{\mathrm{PRF}_{\mathrm{op}}}{2} + \left[\left(\phi - f_k + \frac{\mathrm{PRF}_{\mathrm{op}}}{2}\right) \bmod \mathrm{PRF}_{\mathrm{op}}\right], \tag{4.3}$$

   i.e. with the Doppler of the parcel of sub-band k, whose position is $x(f_k(\phi))$;

2. every channel is corrected with this labelling and the standard reconstruction is run;

3. only sub-band k of the output is kept.

The N_{rx} sub-bands are then put side by side. The cost is N_{rx} SATA passes per channel and N_{rx} reconstructions. In sub-band k the parcel of that sub-band now has the right phase, but the other parcels of the same bin received the phase of the position of parcel k, and the reconstruction lets them leak into it.

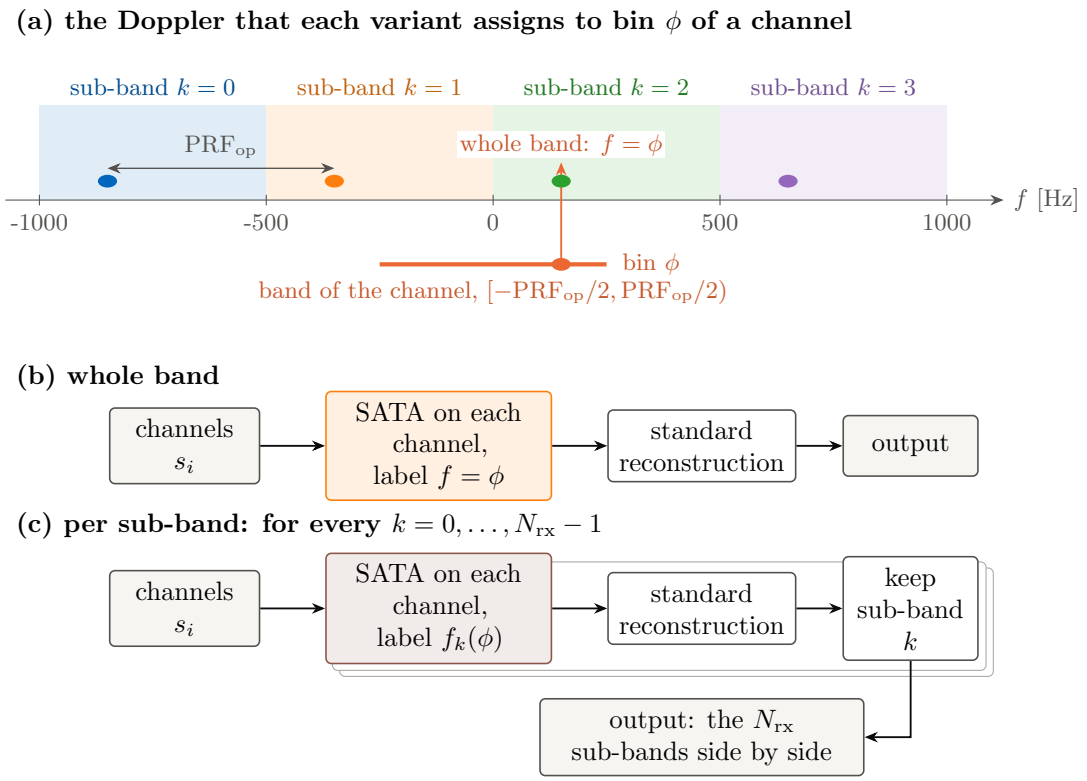


Figure 4.2: The two SATA variants for an aliased channel ($N_{\mathrm{rx}} = 4$, $\mathrm{PRF}_{\mathrm{op}} = 500\,\mathrm{Hz}$). (a) A bin ϕ of the channel holds four parcels; the whole-band variant labels it with $f = \phi$, the per-sub-band variant labels it with its alias $f_k(\phi)$ inside each output sub-band (dots). Through the ruler each label is a ground position, the four labels being X apart. (b) Whole band: one correction per channel, one reconstruction. (c) Per sub-band: one correction and one reconstruction for each k, of which only sub-band k is kept.

Both variants are the identity with a zero map ($< 10^{-15}$) and rotate the line by the right phase with a constant map, and the kernel agrees with the original SATA implementation within 0.44 % rms. Neither, however, can give the parcels of one bin different phases: this is the limit analysed next.

4.2.3 Why it is exact only on special terrain

The channel is aliased: one bin holds the parcels q of the positions $x_q = x_b + \rho f_q$, X apart (Figure 3.1). SATA multiplies the whole bin by a *single* phase, the one of the labelled position x_r. After the correction, parcel q keeps the error

$$\varepsilon_{iq} = \varphi_i(x_q) - \varphi_i(x_r) = -\frac{2\pi\, b_{\mathrm{xt},i}\,\big[h(x_q) - h(x_r)\big]}{\lambda\, r_0 \tan\theta}, \tag{4.4}$$

and the standard reconstruction returns $\hat{\boldsymbol{u}} = \mathbf{L}\boldsymbol{u}$ with $\mathbf{L} = \mathbf{H}^{-1}\mathrm{diag}(\mathrm{e}^{-\mathrm{j}\varepsilon})\mathbf{H}$: the same leakage structure as (3.6), now driven by the *height differences between folded positions*. Hence:

- $\mathbf{L} = \mathbf{I}$, whatever the baselines, only if $h(x) = h(x \pm X) = h(x \pm 2X) = \ldots$: flat terrain, or a terrain periodic with period X;

- for two channels, a $-30\,\mathrm{dB}$ ambiguity tolerates only about $7\,\mathrm{m}$ of height difference between folded positions with $b_{\mathrm{xt}} = 100\,\mathrm{m}$, and $1.4\,\mathrm{m}$ with $500\,\mathrm{m}$;
- in the whole-band variant only the fraction of the target spectrum that falls in the labelled parcel receives the phase of its own position (about 78 %, 52 %, 39 %, 26 % and 20 % for $N_{\mathrm{rx}} = 2, 3, 4, 6, 8$); in the per-sub-band variant the output parcel is right but all the others leak into it.

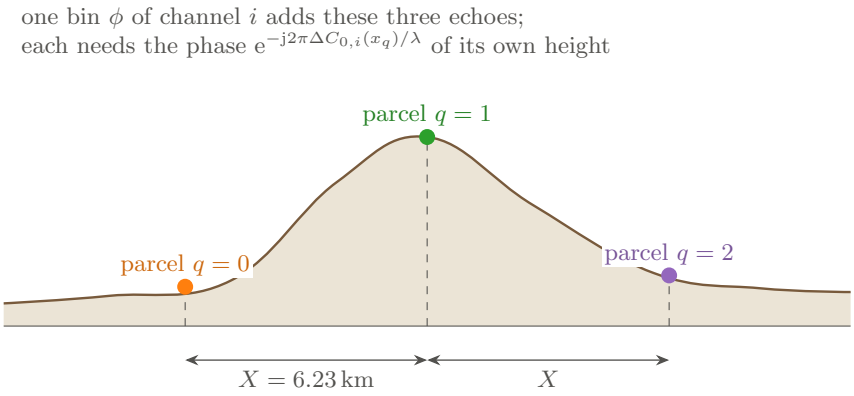


Figure 4.3: The conflict of one bin. Its parcels come from places X apart, with different heights. SATA on the channel applies one phase to the whole bin; the joint reconstruction applies the phase of each place to its own parcel, inside the system of equations.

Correcting each parcel needs the parcels separated, and separating them needs the phases corrected: the two steps cannot be done one after the other. This is what the joint reconstruction of Chapter 5 does in a single system.

4.2.4 Verification

The analysis was checked in three ways (83 automatic checks, all passing):

- **Positive control.** On flat terrain at $500\,\mathrm{m}$, and on the terrain $h = 100 - 100\cos(2\pi x/X)$, whose period is exactly X, SATA whole band, SATA per sub-band and the joint reconstruction coincide (Figure 4.4, left).
- **Negative control.** The same cosine with period $1.5X$: SATA leaves $-30\,\mathrm{dB}$ to $-10\,\mathrm{dB}$ (whole band) and $-24\,\mathrm{dB}$ to $-6\,\mathrm{dB}$ (per sub-band) over the 20 baseline sets, while the joint reconstruction stays within $3\,\mathrm{dB}$ to $6\,\mathrm{dB}$ of the reference in 19 of them ($10\,\mathrm{dB}$ in the last) (Figure 4.4, right).
- **Prediction.** The level left by SATA, computed only from the fold addresses and the DEM, matches the simulation within $1.3\,\mathrm{dB}$ rms (whole band) and $0.7\,\mathrm{dB}$ rms (per sub-band) over 680 runs (Figure 4.5). The residual is structural, not an implementation error.

SATA can even make things worse: a single target at $h = 0$ on a terrain of period $50\,\mathrm{km}$ ($b_{\mathrm{xt}} \le 100\,\mathrm{m}$) reaches $-63.7\,\mathrm{dB}$ without correction, but $-36.3\,\mathrm{dB}$ after SATA whole band and $-27.8\,\mathrm{dB}$ per sub-band, because the neighbouring terrain folded onto the same bins is not at $h = 0$.

A complementary run removes the aliasing (each channel at $\mathrm{PRF}_{\mathrm{op}} = 2\,000\,\mathrm{Hz} > B_{\mathrm{a}}$): there SATA alone brings smooth terrains close to the reference (peak of the worst target within $0.3\,\mathrm{dB}$ to $1.4\,\mathrm{dB}$) and is limited only on steep terrain by its $310\,\mathrm{m}$ cell. The limit discussed above is therefore a consequence of the aliasing, not of SATA itself.

4.2.5 SATA on simulated scenes

Figure 4.4 puts the analysis to the test with five targets on a cosine terrain and four channels ($b_{\mathrm{xt}} \leq 100\,\mathrm{m}$). The left column uses a period of 6.23 km, equal to X: the positions folded onto the same bin (red circles) have the same height, and SATA whole band and per sub-band reach $-44.4\,\mathrm{dB}$, the same level as the joint reconstruction ($-43.9\,\mathrm{dB}$; reference $-48.8\,\mathrm{dB}$). The right column changes only the period, to $1.5X$: the folded positions now have different heights, and SATA leaves $-19.0\,\mathrm{dB}$ (whole band) and $-13.5\,\mathrm{dB}$ (per sub-band), close to the uncorrected $-12.1\,\mathrm{dB}$, while the joint reconstruction does not change ($-44.0\,\mathrm{dB}$).

The residual left by SATA is not an implementation detail. Writing the leakage matrix $\mathbf{L}$ of Section 4.2.3 with the heights of the DEM at the folded positions predicts the level it leaves without any simulation; over the 680 runs of the audit (Section 7.1) the prediction agrees with the simulation within 1.3 dB rms for SATA whole band and 0.7 dB rms per sub-band (Figure 4.5).

The extreme case is a single target at the height of the reference point on a gently varying terrain (Figure 4.6): without correction there is nothing to correct ($-63.7\,\mathrm{dB}$), but SATA gives every bin the phase of the terrain at the labelled position, and the energy of the target receives the phase of other places: $-36.3\,\mathrm{dB}$ whole band and $-27.8\,\mathrm{dB}$ per sub-band. The joint reconstruction leaves the target untouched ($-63.5\,\mathrm{dB}$).

4.2.6 Correction maps that use the target positions

During the project, maps were also built from the simulated targets: the ΔC_0 of each target written at its cell and copied at the folded positions $\pm X$ (*block map with copies*), or the value of the nearest target. With such maps SATA per channel reaches the oracle for isolated targets, because the map then encodes exactly which parcel is occupied. A real scene provides no such information. These maps are kept in the code for reference but are *not* part of any method evaluated in Chapter 7; every method there receives only the system parameters and the DEM, and removing the target list from the configuration given to the methods does not change a single bit of their output.

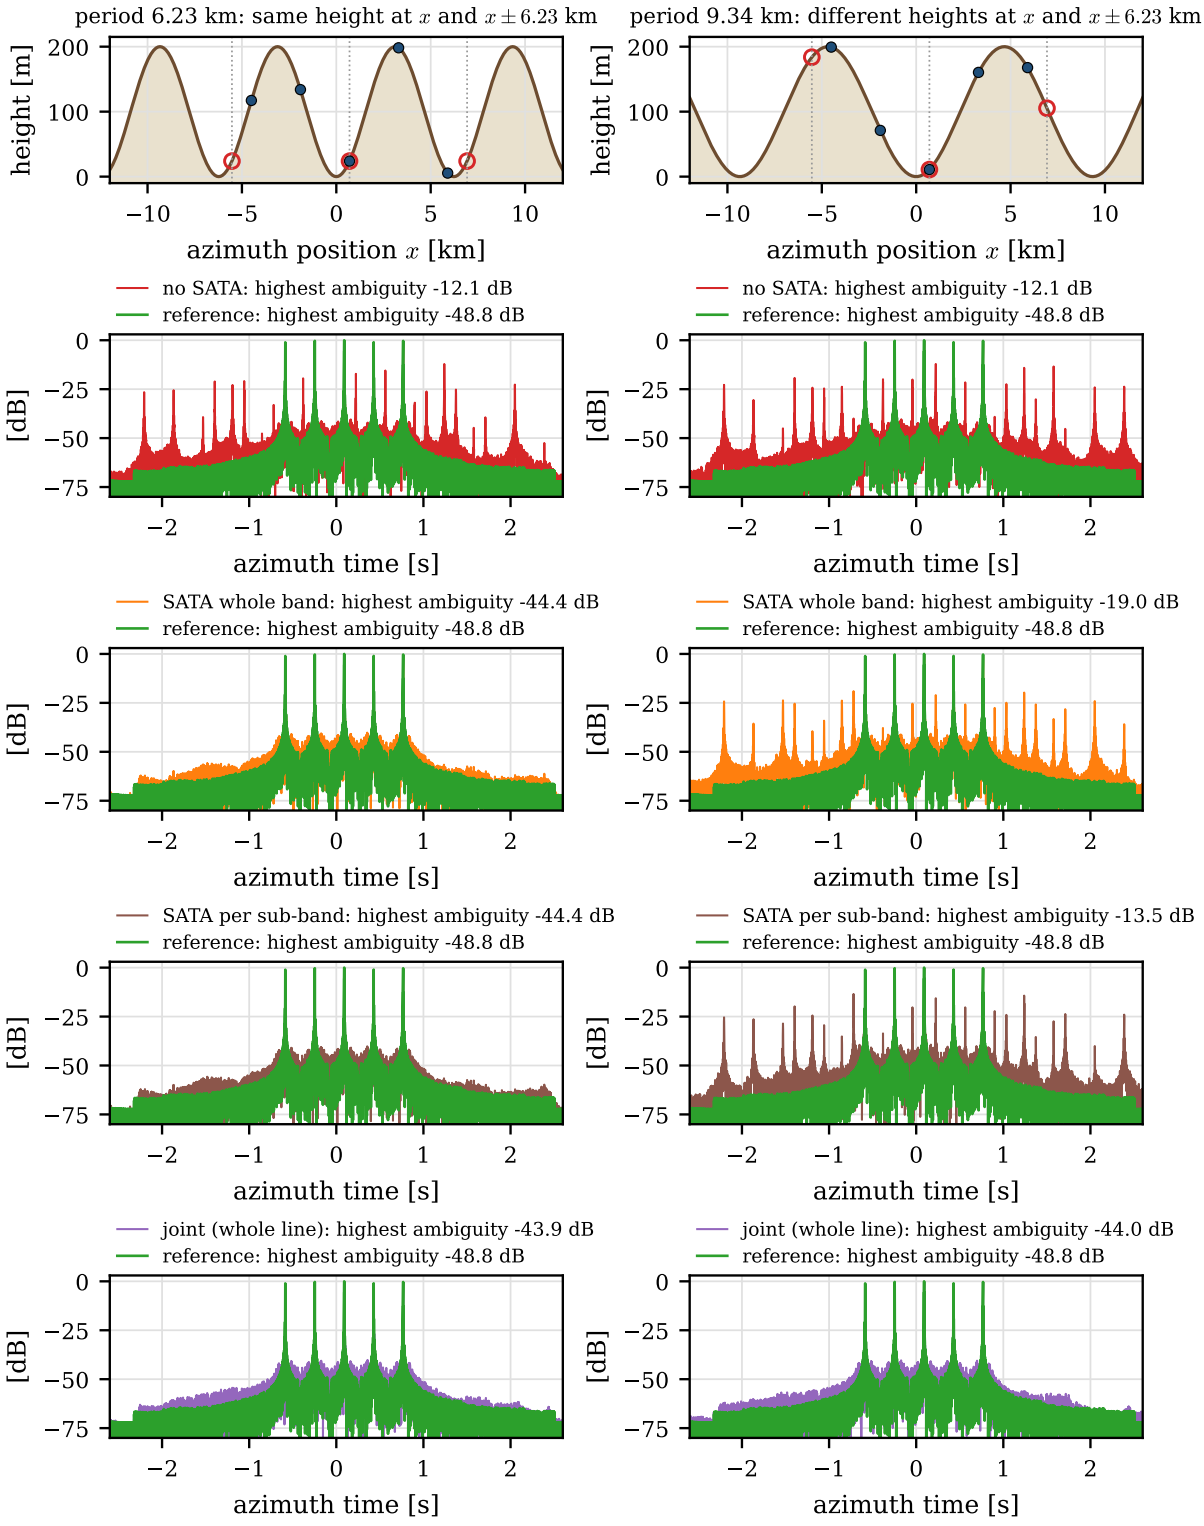


Figure 4.4: Five targets on $h = 100 - 100\cos(2\pi x/P)$, four channels, $b_{\text{xt}} = [0, 100, 39.9, 33.5]$ m. Left: $P = X = 6.23$ km; right: $P = 1.5X$. Top: terrain, targets, and the three positions x, $x \pm X$ folded onto the same channel bin for the middle target (red circles). Below: azimuth impulse responses of the reference and, from top to bottom, of no SATA, SATA whole band, SATA per sub-band and the joint reconstruction over the whole line. Every method reads only the DEM.

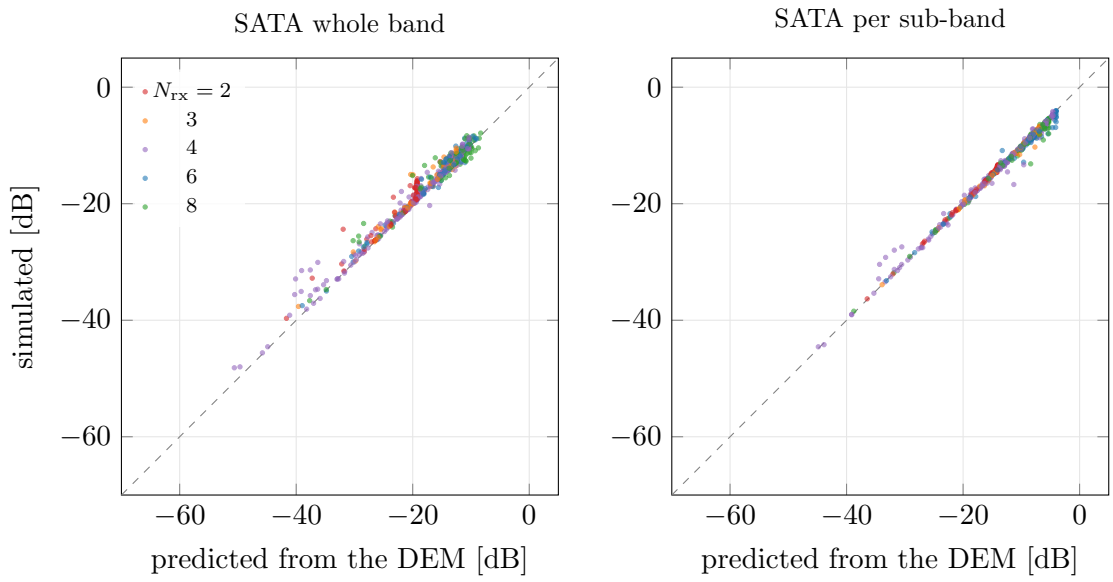


Figure 4.5: Highest ambiguity left by SATA, predicted only from the geometry and the DEM (the leakage of the folded parcels) against the simulation, for the 680 runs of the audit, coloured by the number of channels.

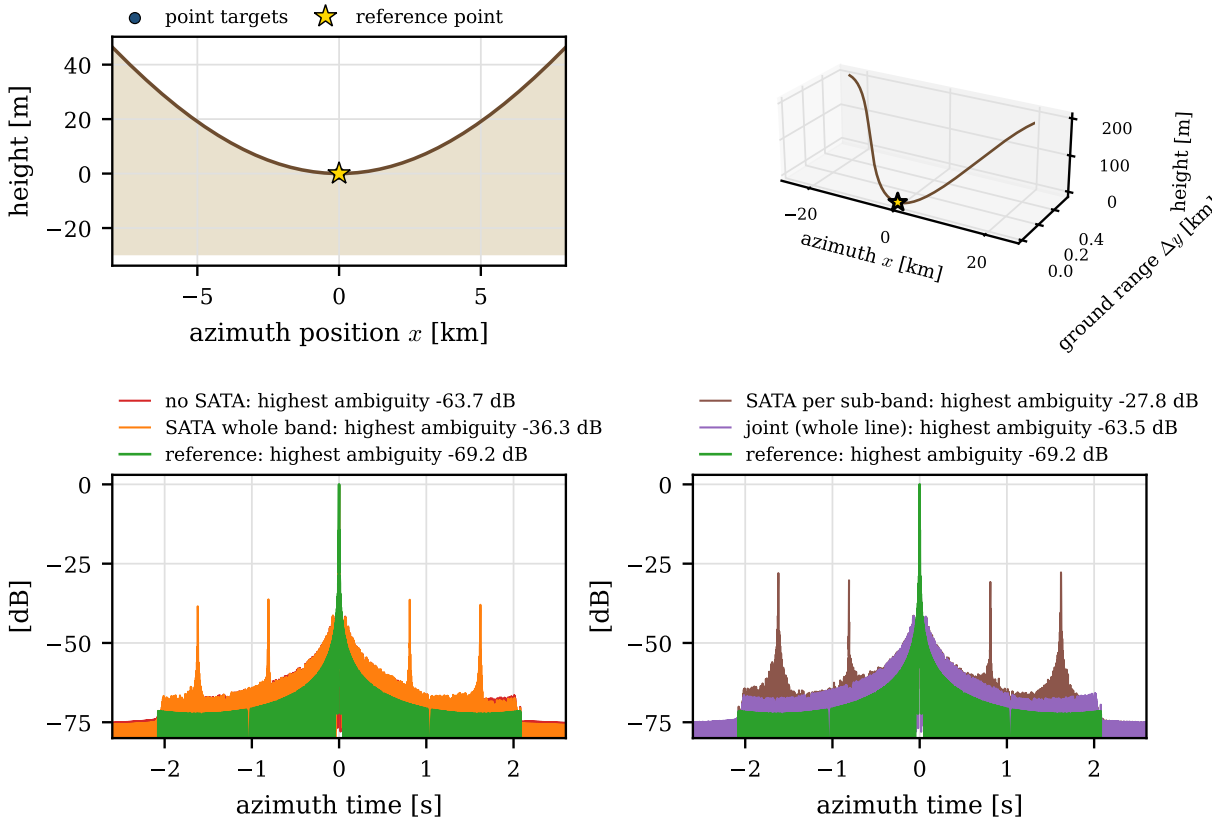


Figure 4.6: One target at the reference point ($x = 0$, $h = 0$) on the terrain $h = 100 - 100\cos(2\pi x/50\,\mathrm{km})$, four channels, $b_{xt} = [0, 100, 39.9, 33.5]$ m. Top: the terrain near the target (left) and along the whole azimuth line (right). Bottom: impulse responses; every method reads only the DEM.

5 Joint topography correction and reconstruction

The analysis of Section 4.2.3 shows that the correct phase (the ΔC_0 read from the DEM) is already known; what fails is applying it per channel before the parcels are separated. The joint reconstruction moves the DEM phase inside the inversion.

5.1 Joint reconstruction in the sub-apertures

With the same sub-apertures as SATA, for every sub-aperture and every bin ϕ of the N_{rx} channel spectra the model (3.2) is written with the phase of each parcel at its own ground position:

$$S_i(\phi) = \sum_q H_i(f_q)\, \exp\!\Big(-\mathrm{j}\frac{2\pi}{\lambda}\Delta C_{0,i}(x_q)\Big)\, U_q, \qquad x_q = x_c + \frac{\lambda r_0}{2v} f_q, \tag{5.1}$$

and solved for the parcels U_q, which are then placed at their true Doppler, transformed back and overlap-added at the full PRF. The steps are summarised in Figure 5.1. With $\Delta C_0 = 0$ this is the standard reconstruction ($\sim 10^{-7}$); with one target it equals the oracle. Its limit is the resolution of the sub-aperture: the DEM is read once per bin of a 40 ms block, i.e. the terrain phase is taken as constant over about ± 600 m of ground. A slope s with a cross-track baseline is equivalent to shifting the phase centre of the channel by $b_{\mathrm{xt}} s/(2\tan\theta)$; when that shift is not small compared with the block, the per-bin model breaks. With $\Delta\varphi = \frac{2\pi}{\lambda}\frac{{b_{\mathrm{xt}}}^{\max} s\,\delta x}{r_0 \tan\theta}$ the phase change inside one SATA cell, the method stays within about 10 dB of the reference for $\Delta\varphi \lesssim 0.5$ rad and fails by 10 to 35 dB above 1 rad.

5.2 Joint reconstruction over the whole line

The final version removes the sub-apertures. The topographic residual of a scatterer is a constant phase over its whole aperture (Chapter 2); in the *focused* full-rate line $g(x)$ it is therefore a pointwise multiplication. The channels are modelled as

$$\boldsymbol{s}_i = \mathcal{D}\Big\{k_i * \big[\, g(x)\, \mathrm{e}^{-\mathrm{j}\varphi_i(x)}\big]\Big\}, \qquad \varphi_i(x) = \frac{2\pi}{\lambda}\Delta C_{0,i}\big(h(x)\big), \tag{5.2}$$

where $\varphi_i(x)$ comes from the DEM only, k_i is the response of channel i to the reference point (the azimuth chirp filtered by the reference-point filter H_i of the standard reconstruction) and $\mathcal{D}$ is the decimation to $\mathrm{PRF}_{\mathrm{op}}$ (Figure 5.2). Writing (5.2) as $\boldsymbol{s} = \mathbf{A}\boldsymbol{g}$, the line is estimated by Tikhonov-regularised least squares,

$$\hat{\boldsymbol{g}} = \arg\min_{\boldsymbol{g}} \sum_i \big\|\boldsymbol{s}_i - \mathbf{A}_i\boldsymbol{g}\big\|^2 + \mu^2\|\boldsymbol{g}\|^2, \tag{5.3}$$

solved with conjugate gradients for least squares (CGLS) [7, 8]. Every iteration costs one application of $\mathbf{A}$ and of its adjoint $\mathbf{A}^{\mathsf{H}}$ (two FFTs per channel each). The regularisation is $\mu = 0.03\,\sigma_{\max}$, with $\sigma_{\max}$ estimated by a few power iterations, and the problem is solved in two passes (iterated Tikhonov: the second pass solves for what the first one left), which turns the filter factor $\sigma^2/(\sigma^2+\mu^2)$ into $1-(1-\sigma^2/(\sigma^2+\mu^2))^2$. Each pass runs up to 60 iterations. The output is re-chirped with the reference response, $U = k * \hat{g}$, so that it has the same form as the standard reconstruction. The complete algorithm is in Figure 5.3.

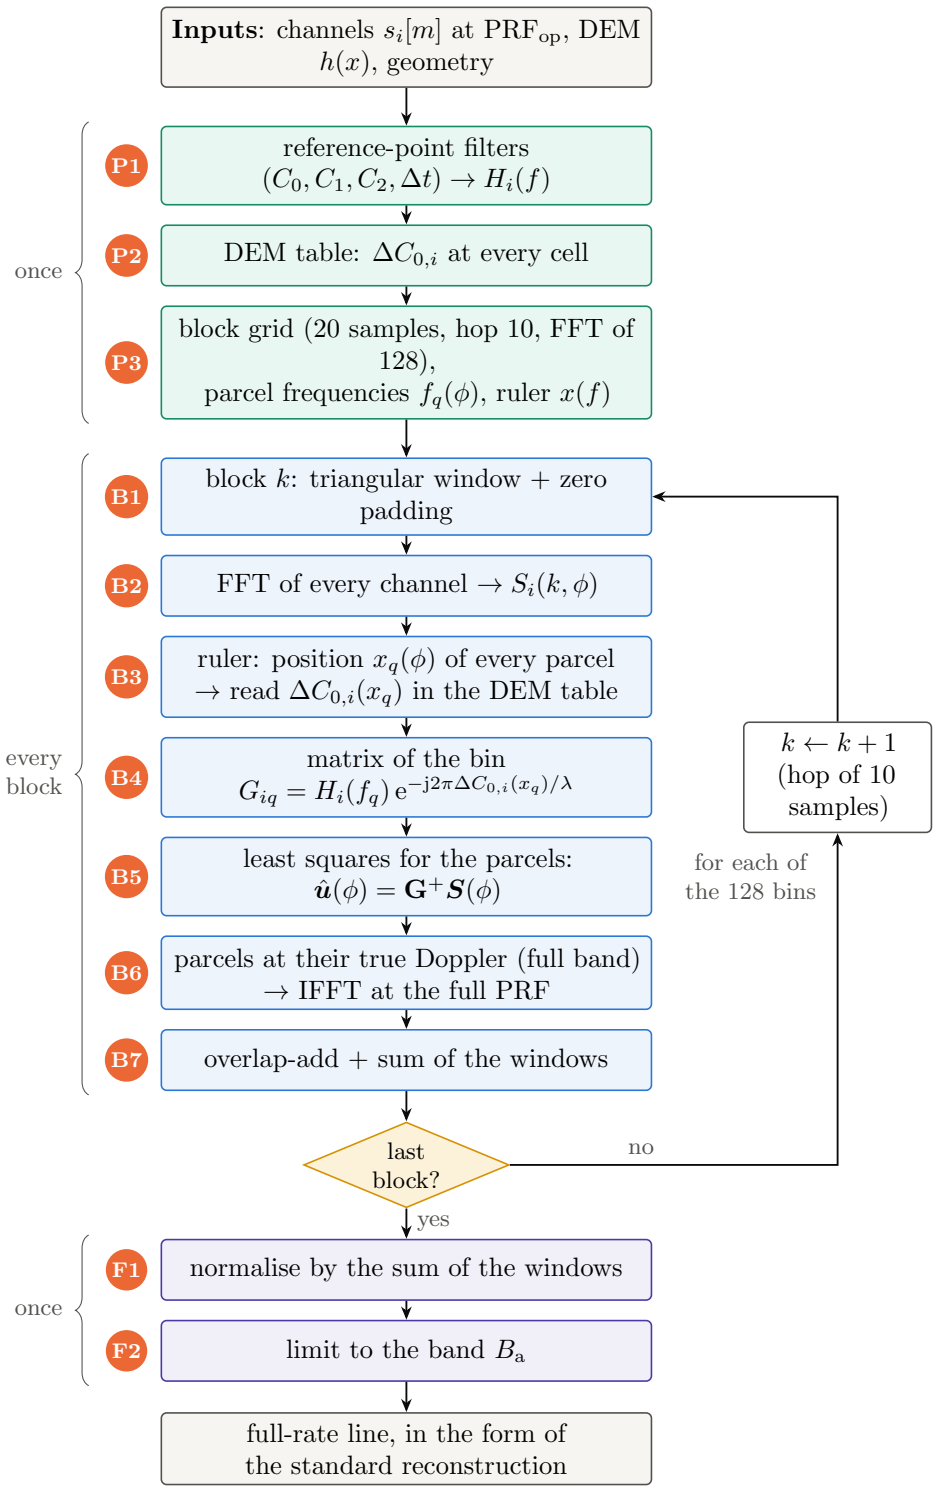


Figure 5.1: The joint reconstruction in the sub-apertures. Green: preparation, done once. Blue: the loop over the blocks of 20 samples; inside each block a small least-squares system is solved for every bin. Violet: finishing. The DEM is read once per bin, at the position the ruler gives to each parcel.

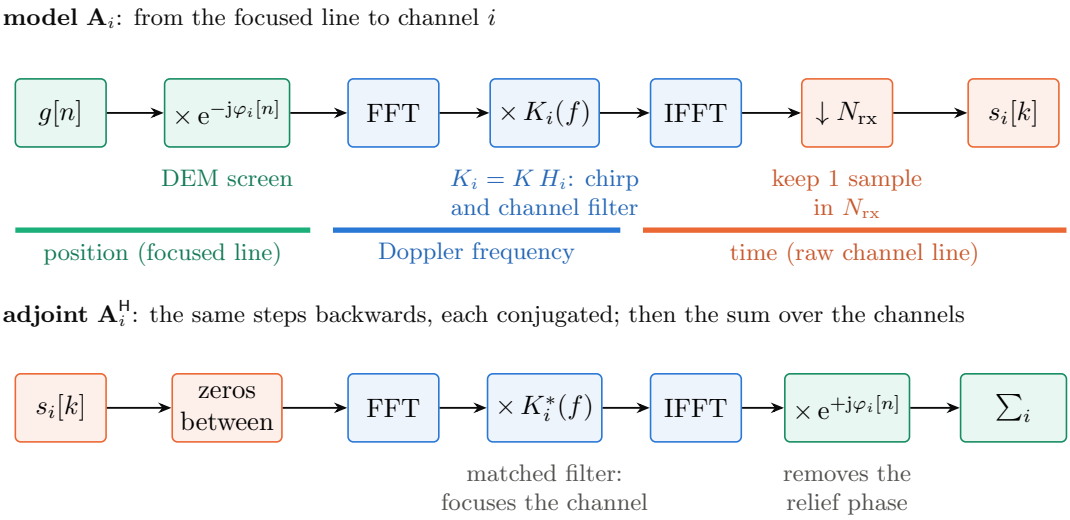


Figure 5.2: The model $\mathbf{A}$ and its adjoint. Each operation is done in the domain where it is simple: the relief phase is a multiplication in *position*, the chirp and the channel filter a multiplication in *Doppler frequency*, the sampling a selection of samples in *time*; the FFTs only move between the domains. The adjoint is the usual SAR processor: it focuses each channel, removes its relief phase and sums the channels.

5.3 Building the model and solving the system

This section builds the model (5.2) from the echo of a single point, one ingredient at a time, writes it as a linear system, shows how that system is solved without ever forming its matrix and explains why the result is the full-rate line $U = K \odot \mathbf{F}\hat{g}$. Unknowns are drawn in red and data in blue.

5.3.1 Building the model

Full-rate line: n is a sample at PRF $= 2\,000$ Hz, $N = 12288$; $\mathcal{F}$ is the DFT; $k[n]$ is the echo of one point (the azimuth chirp) and $K(f) = \mathcal{F}\{k\}$.

Case 1 One target at the reference point, monostatic, sampled at PRF:

$$u[n] = k[n] \xrightarrow{\mathcal{F}} U(f) = K(f).$$

Case 2 Case 1 with the target at $\ell = x/d$, $d = v/\mathrm{PRF} = 3.844$ m:

$$u[n] = k[n-\ell] \xrightarrow{\mathcal{F}} U(f) = K(f)\, e^{-j2\pi f\ell/\mathrm{PRF}}.$$

Case 3 Case 2 with several targets at ℓ_p, with amplitudes a_p:

$$u[n] = \sum_p a_p\, k[n-\ell_p] \xrightarrow{\mathcal{F}} U(f) = \sum_p a_p K(f)\, e^{-j2\pi f\ell_p/\mathrm{PRF}} = K(f) \underbrace{\left[\sum_p a_p\, e^{-j2\pi f\ell_p/\mathrm{PRF}}\right]}_{G(f)=\mathcal{F}\{g\}},$$

$$g[\ell] = \sum_p a_p\, \delta[\ell - \ell_p]. \tag{5.4}$$

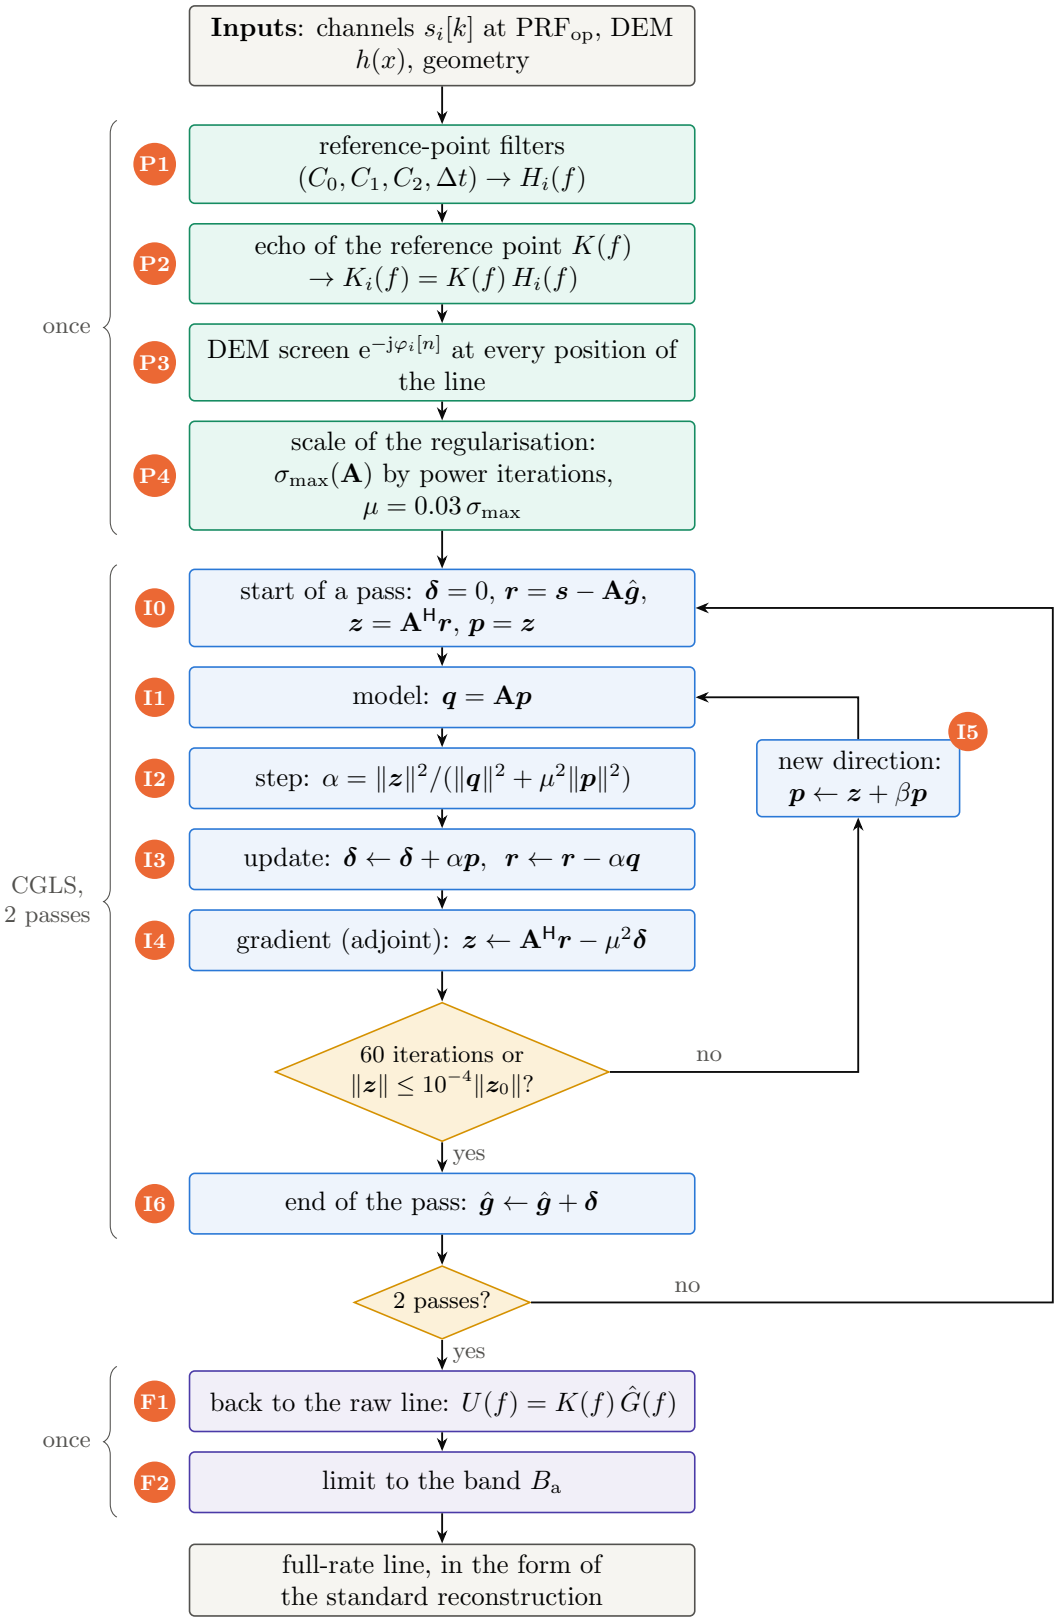
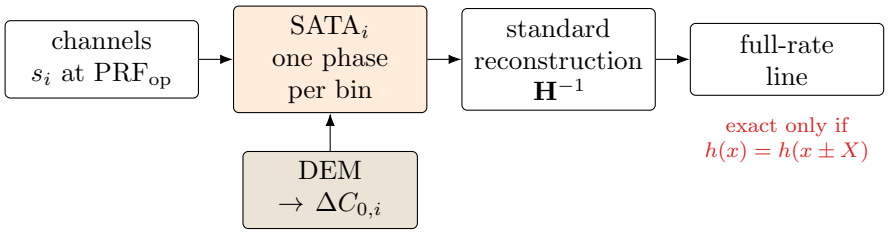


Figure 5.3: The joint reconstruction over the whole line. Green: preparation, done once. Blue: the CGLS iterations, in two passes (iterated Tikhonov); each iteration applies $\mathbf{A}$ and $\mathbf{A}^{\mathsf{H}}$ once (two FFTs per channel each). Violet: back to the raw full-rate line, in the same form as the standard reconstruction.

(a) SATA on each channel, then the standard reconstruction



(b) Joint reconstruction over the whole line (final version)

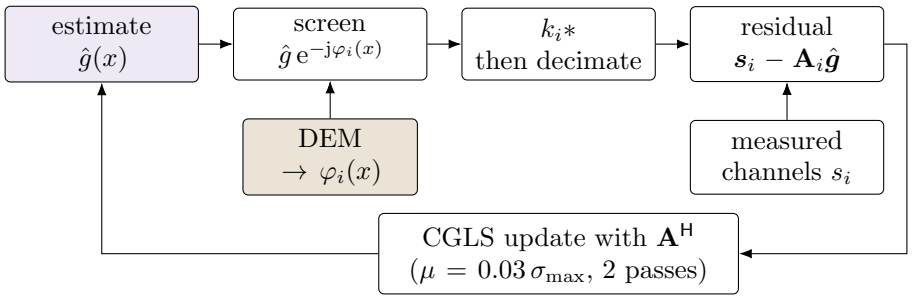


Figure 5.4: Where the DEM phase enters. (a) SATA gives each channel bin a single phase before the parcels are separated. (b) The joint reconstruction puts the phase of every ground position in the model of the channels and inverts model and topography together.

g has one peak per target: it is the focused scene and the unknown of the joint reconstruction.

Case 4 Case 1 with more channels ($i = 0, \dots, N_{\mathrm{rx}} - 1$): each channel is a filter applied to the monostatic echo,

$$K_i(f) := H_i(f)\,K(f), \qquad H_i(f) = \mathrm{e}^{-\mathrm{j}2\pi\left[C_{0,i}/\lambda + (\Delta t_i - C_{1,i})f + C_{2,i}\lambda f^2\right]} \quad (2.5).$$

Case 5 Case 3 + Case 4:

$$S_i(f) = K_i(f)\,G(f).$$

Case 6 Case 5 with relief: in channel i each target gets the phase $\varphi_{i,p} = \frac{2\pi}{\lambda}\Delta C_{0,i}(\ell_p)$ of (2.7),

$$s_i[n] = \sum_p a_p\,\mathrm{e}^{-\mathrm{j}\varphi_{i,p}}\,k_i[n - \ell_p] \;\xrightarrow{\mathcal{F}}\; S_i(f) = K_i(f)\,\underbrace{\Big[\sum_p a_p\,\mathrm{e}^{-\mathrm{j}\varphi_{i,p}}\,\mathrm{e}^{-\mathrm{j}2\pi f\ell_p/\mathrm{PRF}}\Big]}_{\mathcal{F}\{\mathrm{e}^{-\mathrm{j}\varphi_i}\,g\}}$$

$$\Longrightarrow \quad \boxed{S_i(f) = H_i(f)\,K(f)\,\mathcal{F}\{\theta_i \odot g\}(f), \qquad \theta_i[\ell] = \mathrm{e}^{-\mathrm{j}\varphi_i[\ell]}}$$

The phase depends on the *position* ℓ (and on the channel), not on the frequency: it is therefore applied to g, as a pointwise *screen* θ_i, *before* the DFT. This is the reason for choosing the focused line as the unknown: in g the relief of every target is one multiplication, whereas in the raw line the echo of a target spreads over its whole aperture and mixes with the echoes of the terrain X apart, the problem that SATA could not solve (Section 4.2.3).

Case 7 Case 6 recorded by the channel. Case 6 is the full line (2 000 Hz, $N = 12288$); the

channel records one sample in N_{rx} (500 Hz):

$$\boldsymbol{s}_i[k] = s_i^{\mathrm{full}}[N_{\mathrm{rx}}k] \quad \Longrightarrow \quad \boxed{\boldsymbol{s}_i = \underbrace{\mathcal{D}}_{\downarrow N_{\mathrm{rx}}}\ \underbrace{\mathbf{F}^{-1}}_{\mathrm{IFFT}}\ \underbrace{\mathrm{diag}(K_i)}_{\odot K_i}\ \underbrace{\mathbf{F}}_{\mathrm{FFT}}\ \underbrace{\mathrm{diag}(\theta_i)}_{\odot\, \mathrm{e}^{-\mathrm{j}\varphi_i}}\ \boldsymbol{g} = \mathbf{A}_i\boldsymbol{g}} \tag{5.5}$$

with $K_i = H_i \odot K$. Stacking the channels: $\boldsymbol{s} = \mathbf{A}\boldsymbol{g}$, with $\boldsymbol{s} \in \mathbb{C}^{4\times 3072}$ (12288 equations) and $\boldsymbol{g} \in \mathbb{C}^{12288}$ (12288 unknowns). This is (5.2) in matrix form.

5.3.2 What is stored and what is applied

The model and its adjoint are

$$\mathbf{A}_i\boldsymbol{g} = \mathcal{D}\,\mathbf{F}^{-1}\mathrm{diag}(K_i)\,\mathbf{F}\,\mathrm{diag}(\theta_i)\,\boldsymbol{g}, \qquad \mathbf{A}^{\mathsf{H}}\boldsymbol{s} = \sum_i \mathrm{diag}(\theta_i)^{\mathsf{H}}\,\mathbf{F}^{\mathsf{H}}\,\mathrm{diag}(K_i)^{\mathsf{H}}\,(\mathbf{F}^{-1})^{\mathsf{H}}\,\mathcal{D}^{\mathsf{H}}\,\boldsymbol{s}_i, \tag{5.6}$$

with $\mathbf{F}^{\mathsf{H}} = N\mathbf{F}^{-1}$, $(\mathbf{F}^{-1})^{\mathsf{H}} = \mathbf{F}/N$ and $\mathcal{D}^{\mathsf{H}}$ the insertion of $N_{\mathrm{rx}} - 1$ zeros between samples. The adjoint runs the same chain backwards with every factor conjugated, and the channels are summed at the end (Figure 5.2): it focuses each channel with its own filter, removes its relief phase and adds the channels, i.e. it is the usual SAR processor.

$\mathbf{A}$ is a 12288×12288 complex matrix: stored, it would take $12288^2 \times 16\,\mathrm{B} \approx 2.4\,\mathrm{GB}$ and each product would cost about 10^8 operations. Every factor of (5.5), however, is either diagonal or a fast transform, so $\mathbf{A}$ and $\mathbf{A}^{\mathsf{H}}$ are applied as sequences of operations and only the diagonals are kept (Table 5.1): about 1.6 MB and eight FFTs per product with four channels.

Table 5.1: The objects of the model (5.5), and whether each one is ever assembled ($N = 12288$, four channels).

object	meaning	size	assembled?
$\boldsymbol{g}$, $\boldsymbol{s}$	focused scene, channels	N, $4 \times N/N_{\mathrm{rx}}$	vectors
$\mathrm{diag}(\theta_i)$	DEM screen of channel i	$N \times N$, diagonal	only the diagonal
$\mathrm{diag}(K_i)$, $K_i = H_i \odot K$	chirp and channel filter	$N \times N$, diagonal	only the diagonal
$\mathbf{F}$, $\mathbf{F}^{-1}$	DFT	$N \times N$	no: FFT, $\mathcal{O}(N \log N)$
$\mathcal{D}$, $\mathcal{D}^{\mathsf{H}}$	decimation / zero insertion	$N/N_{\mathrm{rx}} \times N$	no: index selection
$\mathbf{A}$, $\mathbf{A}^{\mathsf{H}}$	model and adjoint	$N \times N$	no: chains of the above
$\mathbf{A}^{\mathsf{H}}\mathbf{A} + \mu^2\mathbf{I}$	normal matrix	$N \times N$	never
μ^2	regularisation	scalar	a number

Without relief the screens are constant, $\mathbf{F}\,\mathrm{diag}(\theta_i)\,\mathbf{F}^{-1}$ is a multiple of the identity and the only coupling between Doppler bins is the folding of $\mathcal{D}$, which links a bin f only to $f \pm \mathrm{PRF}_{\mathrm{op}}$, $f \pm 2\mathrm{PRF}_{\mathrm{op}}, \ldots$ Ordered by bin, $\mathbf{A}$ splits into 3072 independent 4×4 blocks (one per bin; rows: channels, columns: parcels). This is exactly the standard reconstruction (3.2), and the only place in the project where a matrix is actually assembled and inverted.

With relief $\mathbf{F}\,\mathrm{diag}(\theta_i)\,\mathbf{F}^{-1}$ is a convolution in Doppler with the spectrum of the screen: it is full and mixes the bins, so the system no longer separates and the whole line is solved at once, iteratively. Figure 5.5 shows one column of $\mathbf{A}^{\mathsf{H}}\mathbf{A}$ for the example of Figure 5.8: on flat terrain a Doppler bin is coupled only with its aliases, more than 40 dB below the bin itself; with the mountain the coupling with the aliases rises to about -22 dB and spreads over ± 7 Hz (some 80 bins).

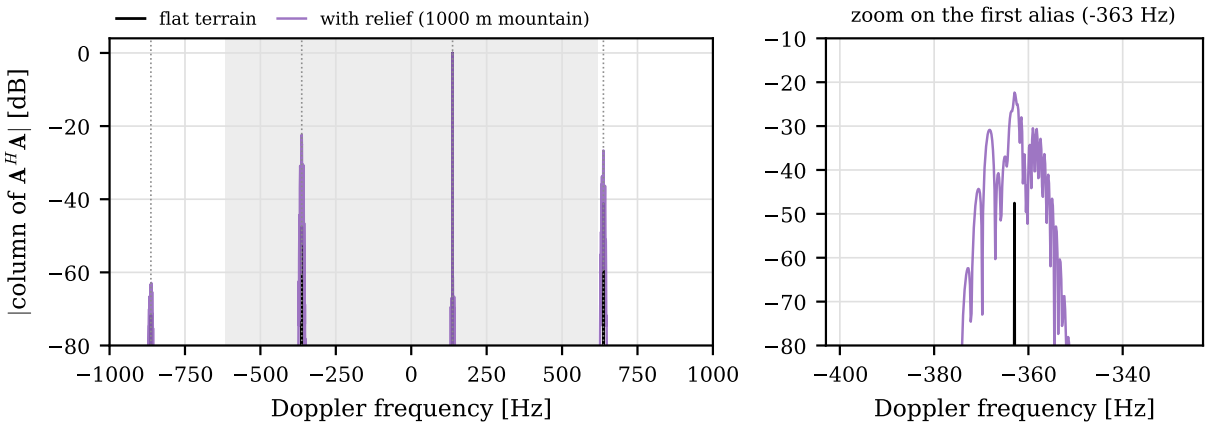


Figure 5.5: One column of $\mathbf{A}^{\mathsf{H}}\mathbf{A}$ in the Doppler domain: the response to a single Doppler bin of g at 137 Hz, for the formation of Figure 5.8, on flat terrain and with the 1 000 m mountain. Dotted lines: the aliases of the bin, $\mathrm{PRF}_{\mathrm{op}} = 500$ Hz apart; grey: the processed band B_{a}. Right: zoom on the first alias.

5.3.3 The problem that is solved

$$\hat{\boldsymbol{g}} = \arg\min_{\boldsymbol{g}} \|\mathbf{A}\boldsymbol{g} - \boldsymbol{s}\|^2 + \mu^2\|\boldsymbol{g}\|^2 \quad \Longleftrightarrow \quad \left(\mathbf{A}^{\mathsf{H}}\mathbf{A} + \mu^2\mathbf{I}\right)\hat{\boldsymbol{g}} = \mathbf{A}^{\mathsf{H}}\boldsymbol{s}, \qquad \mu = 0.03\,\sigma_{\max}(\mathbf{A}). \tag{5.7}$$

The equivalence follows by setting the gradient of the cost to zero, $\mathbf{A}^{\mathsf{H}}(\mathbf{A}\boldsymbol{g} - \boldsymbol{s}) + \mu^2\boldsymbol{g} = \mathbf{0}$. $\sigma_{\max}$ is estimated by the power method, $\boldsymbol{x} \leftarrow \mathbf{A}^{\mathsf{H}}\mathbf{A}\boldsymbol{x}/\|\mathbf{A}^{\mathsf{H}}\mathbf{A}\boldsymbol{x}\|$ (12 iterations), $\sigma_{\max} \approx \sqrt{\|\mathbf{A}^{\mathsf{H}}\mathbf{A}\boldsymbol{x}\|}$, which again uses only $\mathbf{A}$ and $\mathbf{A}^{\mathsf{H}}$.

Why regularise ($\mu^2\|\boldsymbol{g}\|^2$). Writing $\mathbf{A}$ through its singular values, $\mathbf{A}\boldsymbol{v}_k = \sigma_k\boldsymbol{u}_k$, every direction $\boldsymbol{v}_k$ of the scene is recovered with its own weight:

$$\hat{\boldsymbol{g}} = \sum_k \underbrace{\phi_k}_{\text{weight}} (\boldsymbol{u}_k^{\mathsf{H}}\boldsymbol{s})\,\boldsymbol{v}_k, \qquad \phi_k = \begin{cases} 1/\sigma_k & \text{without } \mu \text{ (inverse)}, \\ \dfrac{\sigma_k}{\sigma_k^2 + \mu^2} & \text{with } \mu \text{ (Tikhonov)}. \end{cases} \tag{5.8}$$

A small σ_k is a direction that the channels hardly measure, for example parcels that the channels see with almost the same phases (the conditioning discussed in Section 3.2). Without μ, the factor $1/\sigma_k$ amplifies in those directions the noise and the model errors (the echo shape that the model ignores, the interpolation of the height table, rounding), and $\hat{\boldsymbol{g}}$ explodes. With μ: for $\sigma_k \gg \mu$, $\phi_k \approx 1/\sigma_k$, the same as the inverse; for $\sigma_k \ll \mu$, $\phi_k \approx \sigma_k/\mu^2 \to 0$, the direction is damped. With $\mu = 0.03\,\sigma_{\max}$ only what the channels measure with less than about 3 % of the best direction is damped; the second pass recovers part of what μ removed where the measurement is good.

5.3.4 Conjugate gradients for least squares

$\mathbf{A}^{\mathsf{H}}\mathbf{A} + \mu^2\mathbf{I}$ is Hermitian and positive definite, so (5.7) can be solved by conjugate gradients. CGLS does it with products by $\mathbf{A}$ and $\mathbf{A}^{\mathsf{H}}$ only, never forming the normal matrix (which would also square the condition number in finite precision). The loop is in Figure 5.3:

- each iteration costs one model ($\boldsymbol{q} = \mathbf{A}\boldsymbol{p}$) and one adjoint ($\boldsymbol{z} = \mathbf{A}^{\mathsf{H}}\boldsymbol{r} - \mu^2\boldsymbol{\delta}$); $\alpha = \|\boldsymbol{z}\|^2/(\|\boldsymbol{q}\|^2 + \mu^2\|\boldsymbol{p}\|^2)$ is the optimal step along the direction $\boldsymbol{p}$, and $\boldsymbol{p} \leftarrow \boldsymbol{z} + \beta\boldsymbol{p}$, $\beta = \|\boldsymbol{z}_{\mathrm{new}}\|^2/\|\boldsymbol{z}\|^2$, is conjugate to the previous directions, so no direction is searched twice;

- a pass stops when $\|\boldsymbol{z}\| \leq 10^{-4}\|\boldsymbol{z}_0\|$ or after 60 iterations, and adds its result to $\hat{\boldsymbol{g}}$;
- two passes: the second solves the same problem for the residual left by the first and adds its solution to $\hat{\boldsymbol{g}}$ (iterated Tikhonov). The weight becomes $\left[1-\left(\mu^2/(\sigma_k^2+\mu^2)\right)^2\right]/\sigma_k$, closer to $1/\sigma_k$ where the channels measure well and still zero where they do not;
- on the example of Figure 5.8: 120 iterations in total, with $-40.3\,\mathrm{dB}$ of the channels left unexplained.

5.3.5 The answer: $U = K \odot \mathbf{F}\hat{\boldsymbol{g}}$

What is wanted. The line that an ideal *monostatic* radar, at the transmitter, would record at PRF $= 2\,000\,\mathrm{Hz}$, without aliasing. In the simulation this is the reference. $\hat{\boldsymbol{g}}$ is the scene; U is that scene recorded again by the ideal radar.

Where it comes from. It is Case 3, (5.4), applied to the scene that was found:

$$U(f) = \underbrace{K(f)}_{\text{echo of one point}} \underbrace{\hat{G}(f)}_{\mathcal{F}\{\hat{g}\}} \iff u[n] = \sum_{\ell} \hat{g}[\ell]\, k[n-\ell], \tag{5.9}$$

i.e. one chirp at each position of the scene, weighted by $\hat{g}$ and summed.

Why without the screen. The screen is the phase of channel i *relative to the monostatic* path; for the monostatic radar itself it is one ($b_{\mathrm{at}} = b_{\mathrm{xt}} = 0 \Rightarrow C_0 = 0 \Rightarrow \Delta C_0 = 0 \Rightarrow \theta \equiv 1$). The effects of the topography are therefore split as in Table 5.2: the monostatic phase of each target is common to all channels and is part of its amplitude a_p, so it is estimated inside $\hat{\boldsymbol{g}}$ and reaches U, exactly as in the reference; the channel-dependent phase lives in the model $\mathbf{A}$ and is used only to explain the channels. U thus contains what the reference contains, the scene with its monostatic phases, and none of the phases that differ between channels, which are the ones that create the ambiguities.

Table 5.2: How each effect of the topography enters the joint model and the answer U.

effect of the topography	type	in the model	in U?
monostatic range phase, $\mathrm{e}^{-\mathrm{j}4\pi r_p/\lambda}$	constant per target, same in every channel	in a_p (inside $\boldsymbol{g}$)	yes, through $\hat{\boldsymbol{g}}$
extra phase of channel i, $\varphi_i = 2\pi\Delta C_{0,i}/\lambda$	constant per target, differs per channel	screen θ_i (inside $\mathbf{A}$)	no ($\theta \equiv 1$)
echo shape ($K_a \propto 1/r$; C_1, C_2 with h)	not constant	ignored ($\Delta r/r \approx 0.1\,\%$)	—

Why back to the raw line, and not stop at $\hat{\boldsymbol{g}}$.

- *Fair comparison:* the reference, the standard reconstruction, SATA and the joint reconstruction all return the raw line of 12288 samples at $2\,000\,\mathrm{Hz}$; the focusing that follows is the same for all of them.
- *Band:* outside $|f| \leq B_{\mathrm{a}}/2$ the data do not determine $\boldsymbol{g}$; the multiplication by K and the band cut discard that part.
- *SAR chain:* the reconstruction is a pre-processing step (aliased channels in, one clean channel out); the SAR processor expects raw data.

In practice $U = K \odot \mathbf{F}\hat{\boldsymbol{g}}$, then $U(f) = 0$ for $|f| > B_a/2$ ($B_a = 1\,234.8$ Hz), and the line is $\mathbf{F}^{-1}U$, in the same format as the standard reconstruction.

Afterwards (the same for every method). The focusing is image $= \mathbf{F}^{-1}\{K^*U\} = \mathbf{F}^{-1}\{|K|^2\hat{G}\}$, i.e. $\hat{G}$ with a smooth window inside the band: one peak per target (-3 dB width ≈ 5.8 m). The metric of Section 3.3 is then the highest ghost at the folds relative to the peak of the reference. Figure 5.6 summarises the chain.

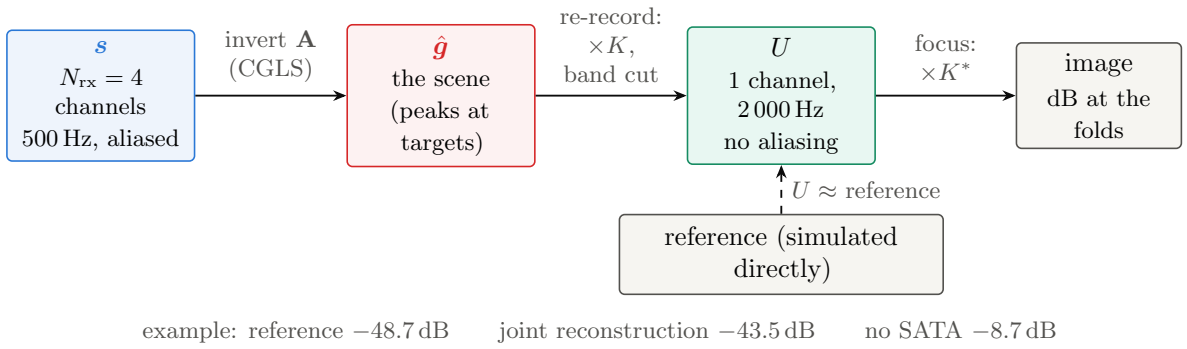


Figure 5.6: From the aliased channels to the image. The joint reconstruction inverts the model to find the scene $\hat{\boldsymbol{g}}$, records it again with the ideal monostatic radar ($\times K$ and band cut) to obtain U, which is compared with the reference simulated directly; the focusing is the same for every method. Values: example of Figure 5.8.

5.4 Properties

- **Consistency.** With a zero DEM the method reproduces the standard reconstruction within 0.1 % rms.
- **Information used.** Only the channels, the system geometry and the DEM. The positions of the targets are never used (Section 4.2.6).
- **Self-test of the DEM.** The part of the channels that the model does not explain is -41 dB to -39 dB with the right DEM and -20 dB to -7 dB with a wrong one (mirrored or shifted by 1 km), which makes it a natural alarm for DEM errors.
- **Cost.** About 3.3 s per line with four channels in the 1-D simulator (median over 680 runs), against 0.15 s for SATA whole band.
- **Convergence.** Figure 5.7 follows the iterations on the example of Figure 5.8. The first iteration is the adjoint, i.e. the channels focused, corrected and summed (-10.2 dB, close to no SATA); about ten iterations bring the ambiguities below -39 dB, and the second pass gains another 1 dB (-43.5 dB after 2×60 iterations). The unexplained part of the channels falls to -38 dB in the same way.
- **Ill-conditioned configurations.** Where the baselines do not separate the parcels (two channels, equally spaced baselines, certain heights with long baselines) the information is not in the data; the regularisation then leaves the ambiguity instead of amplifying noise.

Figure 5.8 shows the example used throughout the development: five targets on a 1 000 m Gaussian mountain with $b_{xt} = [0, 134.9, 20.5, 8.3]$ m. The highest ambiguity goes from -8.7 dB without correction to -43.5 dB with the joint reconstruction; the reference is -48.7 dB, and -43.9 dB once it is limited to the processed band B_a, as every reconstruction is.

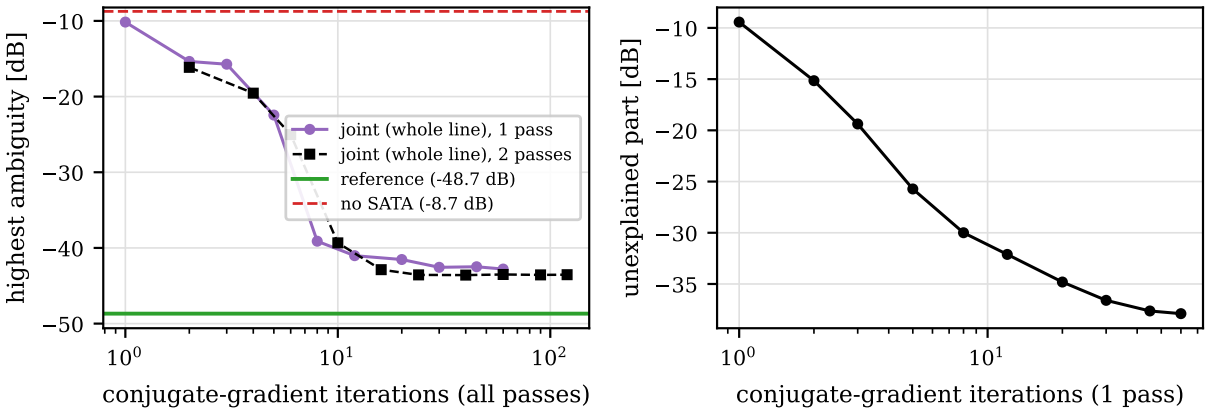


Figure 5.7: Joint reconstruction over the whole line on the example of Figure 5.8 (five targets on a 1 000 m mountain, four channels). Left: highest ambiguity against the total number of conjugate-gradient iterations, with one pass and with two passes, against the reference and no SATA. Right: part of the channels that the model does not explain, $\|\boldsymbol{s} - \mathbf{A}\hat{\boldsymbol{g}}\|/\|\boldsymbol{s}\|$, during the first pass.

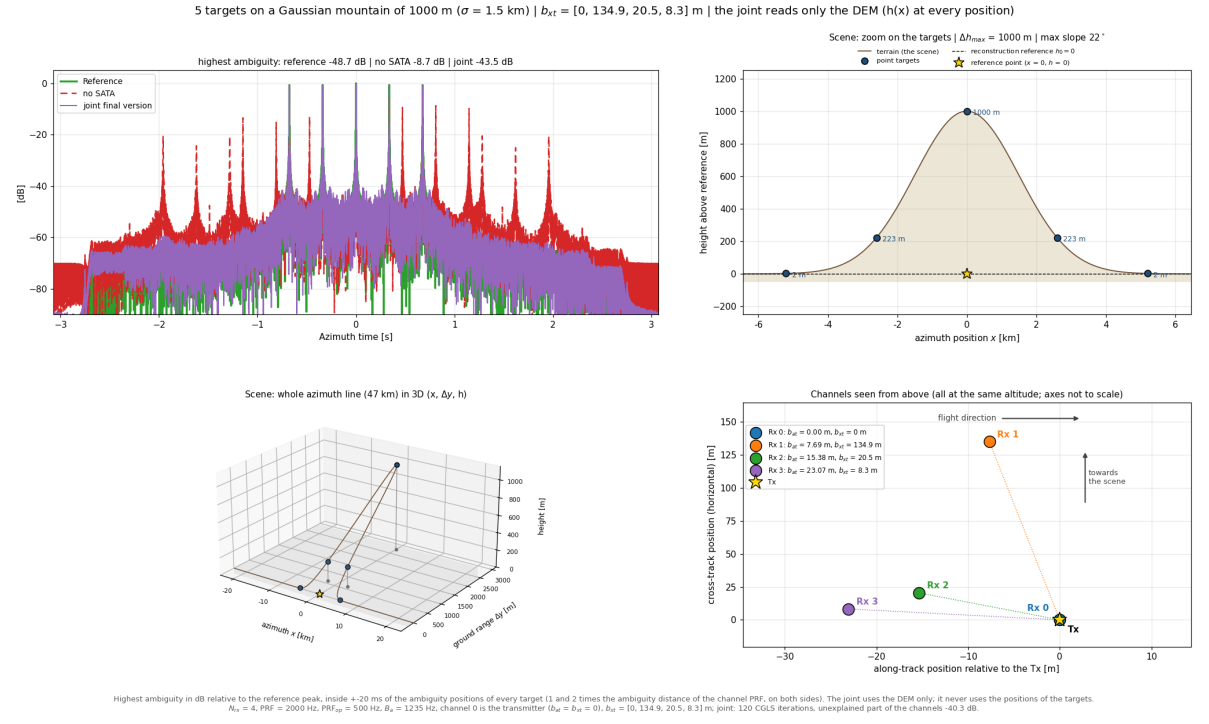


Figure 5.8: Five targets on a 1 000 m Gaussian mountain, four channels. Top left: azimuth impulse responses of the reference, of the standard reconstruction (no SATA) and of the joint reconstruction over the whole line. Right: the scene and the formation. The joint reconstruction reads only the DEM.

6 Extension to two dimensions

The 1-D study keeps every target at the slant range of the reference point. In a real scene the terrain varies in range as well, the residual $\Delta C_{0,i}$ depends on both coordinates, the echoes migrate in range and the range envelope of each channel is shifted by half its path difference. A second simulator handles 3-D scenes with along- and cross-track baselines and range-compressed data, and the joint reconstruction is extended to it.

6.1 Model

The unknown is the focused image $g[n, m]$ (Figure 6.2) on azimuth positions x_n (full PRF) and slant ranges of closest approach r_m. The range-compressed data of channel i are modelled as

$$\boldsymbol{s}_i = \mathcal{D}\Big\{ K_i \cdot \mathcal{F}_r\big[\, Q_i \cdot \mathcal{F}_a(S_i\, g)\big]\Big\} \;+\; (\text{envelope term}), \tag{6.1}$$

with

- S_i, the terrain screen at every pixel, $\exp\big(-\mathrm{j}2\pi\Delta C_{0,i}(x, r)/\lambda\big)$, computed from the DEM in radar geometry (the ground range of each pixel is found by a fixed-point iteration, Figure 6.1, which converges for range slopes below $\tan\theta$, i.e. without layover);
- $Q_i(f; r)$, what changes with the range of closest approach in the range-Doppler domain: the azimuth chirp relative to the reference range and the differences of the coefficients $C_0, C_1, C_2, \Delta t$ of the channel;
- $K_i(f, f_r)$, the 2-D kernel: the 2-D FFT of the simulated echo of the reference point in channel i at the full PRF (chirp, range migration, window of the pair, co-registration by $C_{0,i}/2$);
- the envelope term, the first-order shift of the range envelope by $\frac{1}{2}[\Delta C_{0,i} + C_{0,i}(r) - C_{0,i}(r_{\mathrm{ref}})]$, valid while this shift is small compared with the range resolution.

The image is kept on a range grid twice as fine as the data, because a phase that varies with range shifts the range spectrum. The adjoint passes the adjoint test to 2.6×10^{-7}; the solver is the same CGLS as in 1-D ($\mu = 0.03\,\sigma_{\max}$, two passes of up to 60 iterations).

6.2 The common range band

A cross-track baseline shifts the range spectrum seen by the channel, the spectral shift known from interferometry:

$$\nu_i = \frac{c\, b_{\mathrm{xt},i} \cos\theta}{2\lambda\, r \tan(\theta - \alpha)}, \tag{6.2}$$

with α the range slope of the terrain; on flat terrain this is 0.202 MHz per 100 m of baseline. Range frequencies that only some channels observe are not determined by the data (fewer equations than folded parcels), and the joint reconstruction leaves the full ambiguity there. In a control case (one target, flat terrain, $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m, exact model) the full band reaches -40.5 dB, whereas the band common to all four channels reaches -66.1 dB, equal to the reference in the same band. The output is therefore limited, pixel by pixel, to the range frequencies that

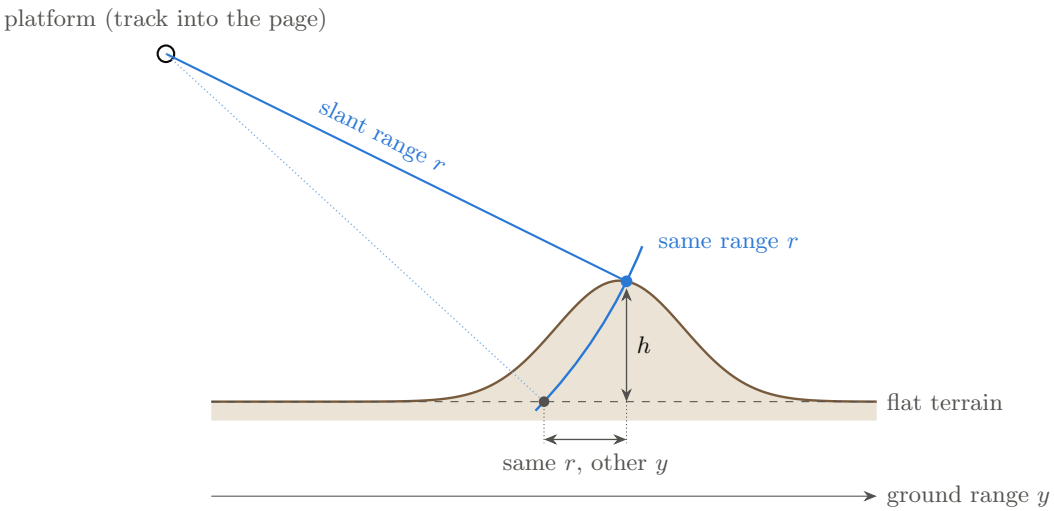
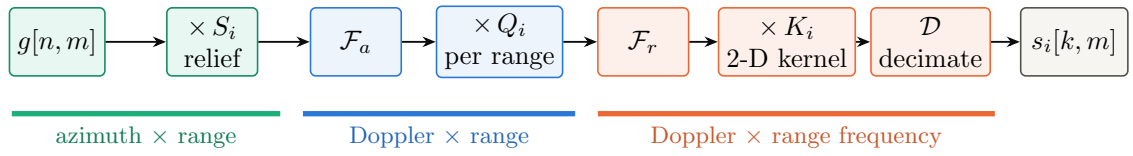


Figure 6.1: The DEM in radar geometry: a pixel at slant range r lies on the terrain at another ground range than on flat terrain. The ground range is found by the fixed-point iteration $y \leftarrow \sqrt{r^2 - [H - h(x,y)]^2}$, which contracts while the range slope stays below $\tan\theta$.

model: from the focused image to the range-compressed data of channel i



adjoint: repeat the spectrum, 2-D matched filter, remove the relief, sum the channels

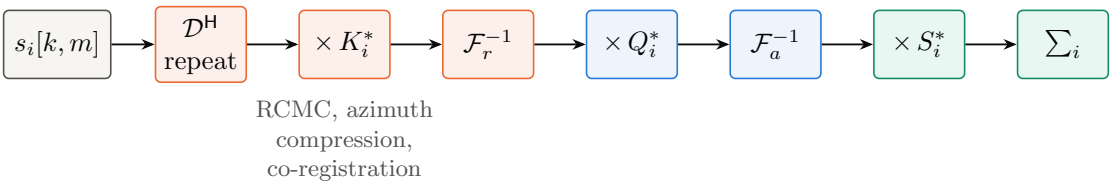


Figure 6.2: The model of the 2-D joint reconstruction and its adjoint (the envelope term is omitted). The relief acts per pixel, the range-dependent focusing per range line in the range-Doppler domain, and the 2-D kernel of the channel in the two-dimensional frequency domain.

every channel observes there, and the reference and the uncorrected reconstruction are compared in the same band. The price is range resolution, a fraction $b_{\mathrm{xt}}/{b_{\mathrm{xt}}}^{\mathrm{crit}}$ of the band; with the 5 MHz bandwidth of the sweeps the median pixel keeps 93 % of the band and the worst pixel 45 %. A wider system bandwidth (e.g. 50 MHz) would reduce this loss about tenfold.

Figure 6.3 shows three 3-D scenes of the 2-D sweep, each with 49 targets spread over 5.4 km of azimuth and 2.4 km of ground range: an oblique ramp and a U-shaped valley with four channels and ${b_{\mathrm{xt}}}^{\max} = 300\,\mathrm{m}$, and the isolated 1 000 m mountain with six channels and ${b_{\mathrm{xt}}}^{\max} = 500\,\mathrm{m}$. Without correction the highest ambiguities are $-10.6\,\mathrm{dB}$, $-12.3\,\mathrm{dB}$ and $-8.6\,\mathrm{dB}$; the joint reconstruction, reading only the DEM, brings them to within 1 dB of the reference in all three. Not every case is this good: the same mountain with four channels and $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]\,\mathrm{m}$ keeps $-25.1\,\mathrm{dB}$ (reference $-44.4\,\mathrm{dB}$), one of the hard configurations of Chapter 7, while with the baselines $[0, 100, 39.9, 33.5]\,\mathrm{m}$ it reaches the reference exactly.

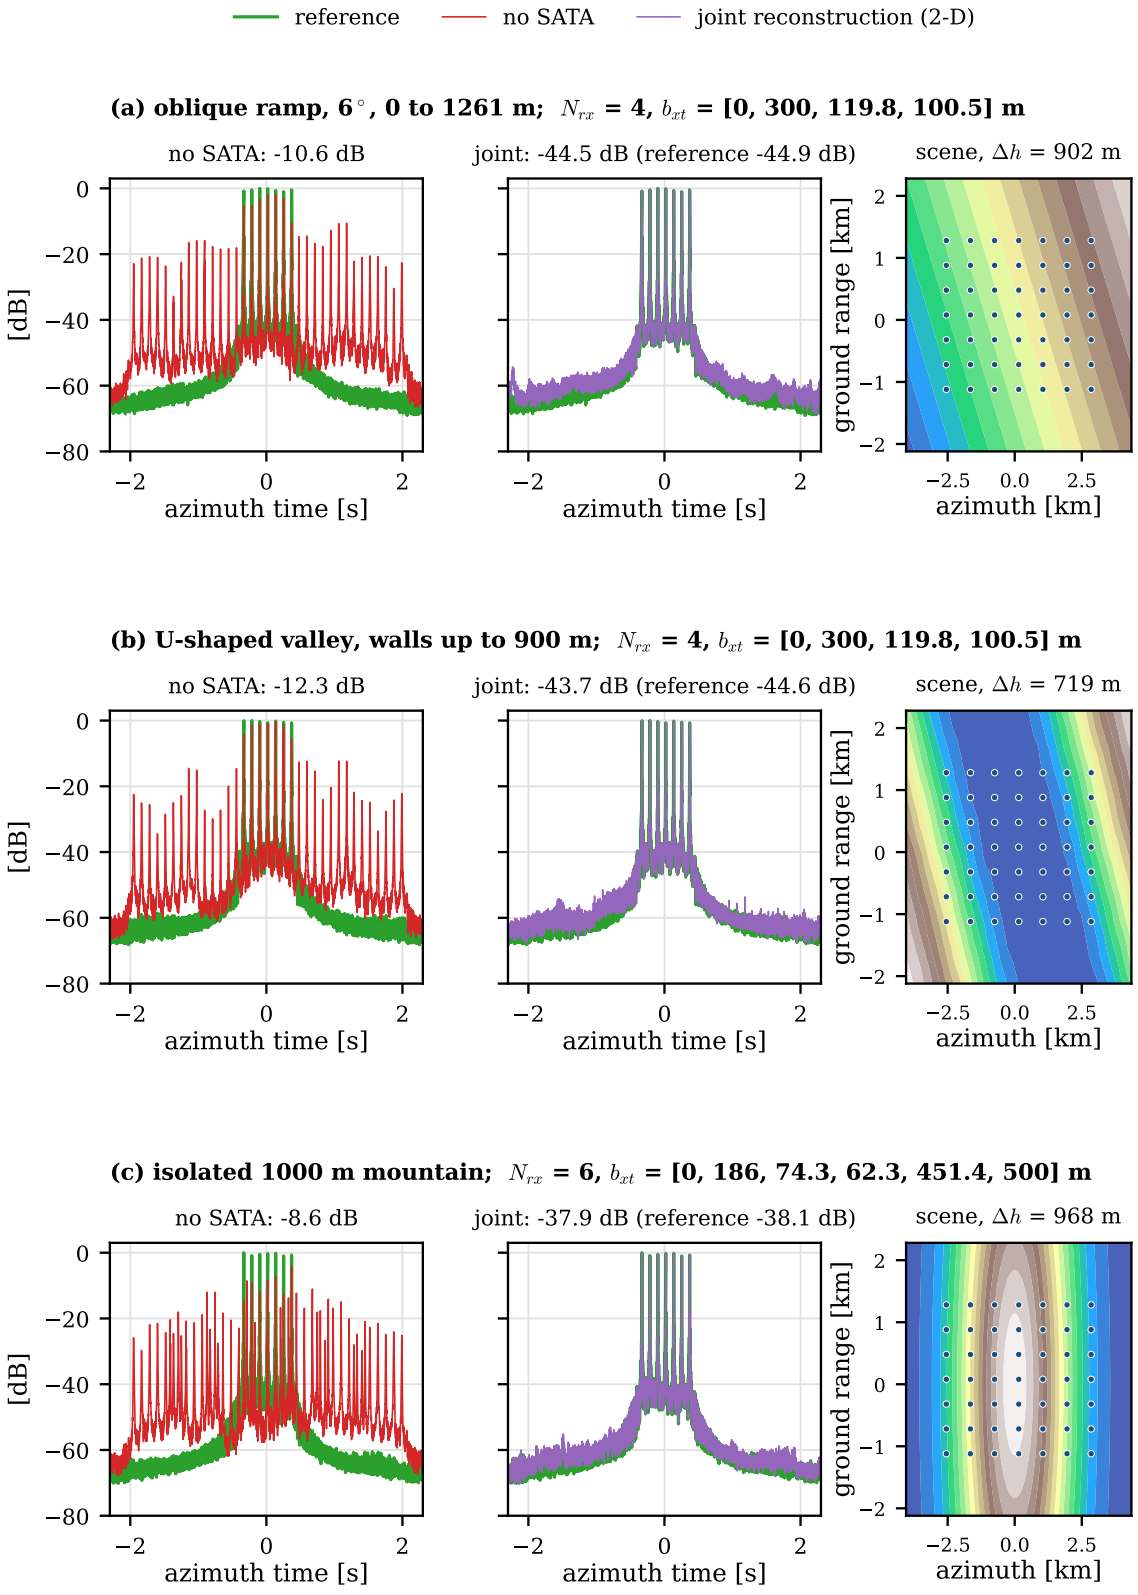


Figure 6.3: Azimuth impulse responses of three scenes of the 2-D sweep (49 targets each, every target focused with its own 2-D matched filter, envelope over the targets; all curves in the same azimuth and common range band). Left: reference and no correction; centre: reference and the 2-D joint reconstruction; right: the terrain seen from above with the targets. The joint reconstruction reads only the DEM. (a) Oblique ramp of 6°, 35° off the flight direction; (b) glacial valley with walls up to 900 m, axis 60° off the flight direction; (c) isolated 1 000 m Gaussian mountain.

7 Results

All numbers in this chapter are the highest ambiguity of Section 3.3, in dB relative to the peak of the reference. Every method receives only the system parameters and the DEM. The scenes cover along-track ramps (3° to 20°), sines and cosines of several periods, and realistic terrains (hills, mountains, escarpments, plateaus, fractal alpine relief, glacial valleys, ridges, canyons, volcanoes), each with one, a few and many point targets (up to 49). The baselines are $N_{\mathrm{rx}} = 2, 3, 4, 6, 8$ with ${b_{\mathrm{xt}}}^{\max} = 100\,\mathrm{m}$, 300 m and 500 m (random, the largest exactly ${b_{\mathrm{xt}}}^{\max}$), plus five fixed four-channel sets: two random ones (up to 100 m and up to 500 m), two equally spaced ones ($[0, 10, 20, 30]$ and $[0, 50, 100, 150]$ m) and one with baselines on both sides of the transmitter ($[0, -50, 50, 100]$ m). Channel 0 is always the transmitter.

7.1 All methods with a DEM only

Table 7.1 and Figure 7.1 compare the methods on 680 runs (34 scenes × 20 baseline sets; median reference −48.7 dB).

Table 7.1: All methods with the DEM only (680 runs).

Method	Median highest ambiguity [dB]	Runs within 10 dB of the reference [%]
No correction (standard reconstruction)	−8.7	3.5
SATA whole band, then reconstruction	−15.4	2.5
SATA per sub-band	−10.8	1.2
Joint reconstruction in the sub-apertures	−30.9	36.3
Joint reconstruction over the whole line	**−41.4**	**81.9**

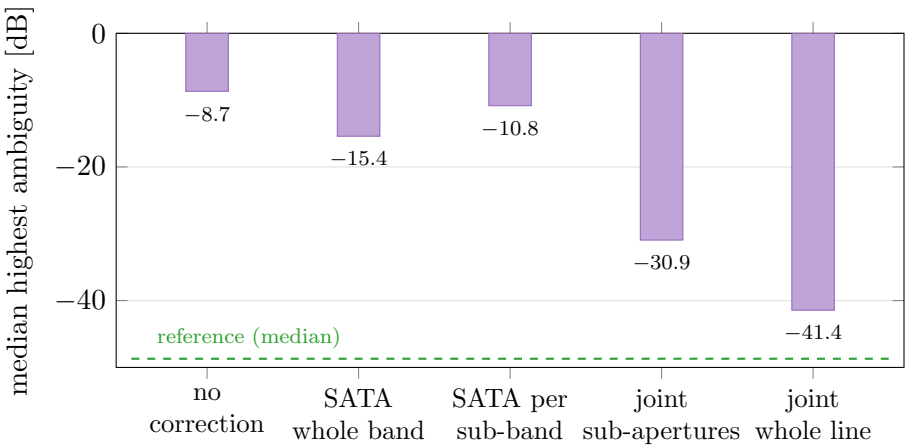


Figure 7.1: Median highest ambiguity of every method over the 680 runs of the audit. SATA per channel stays close to the uncorrected reconstruction; the joint reconstruction over the whole line gets within a few dB of the reference.

Head to head, the joint reconstruction over the whole line is better than SATA whole band in all 680 runs (median improvement 25.7 dB) and better than SATA per sub-band in all of them (median 30.5 dB). SATA per channel improves on the uncorrected reconstruction by a median of

only 5.5 dB (whole band) and 0.2 dB (per sub-band), consistently with Section 4.2.3.

A distributed scene (about 11 900 scatterers, oracle −34.2 dB, measured as error-to-signal ratio) confirms the ranking: with $b_{\mathrm{xt}} \sim U(0, 100)$ m SATA per channel gives −22.5 dB to −8.7 dB and the whole-line joint −34.1 dB to −33.2 dB; with $U(0, 500)$ m, −13 dB to −6 dB and −33.6 dB to −31.4 dB.

7.2 Why the joint reconstruction is the final method

Figure 7.2 compares all the methods on one realistic scene: five targets on random mountains (relief 1 000 m), four channels with ${b_{\mathrm{xt}}}^{\max} = 300$ m. No SATA leaves −4.6 dB; SATA whole band −14.6 dB; SATA per sub-band −8.6 dB; the joint reconstruction over the whole line −43.6 dB, 5 dB from the reference (−48.7 dB). Figure 7.3 shows that this is not a particular case: over the 680 runs of the audit the joint reconstruction is better than SATA in every run, by a median of 25.7 dB (whole band) and 30.5 dB (per sub-band).

The reasons, collected from the previous chapters:

1. **SATA has the right phase in the wrong place.** SATA reads the correct correction, the ΔC_0 of the DEM. But on an aliased channel each bin is the sum of parcels from positions $X \approx 6.2$ km apart, and a per-channel correction can give that sum only one phase (Figure 4.3). The standard reconstruction that follows assumes that every channel is now exact, so the error of the other parcels becomes leakage between the parcels. SATA is exact only when the terrain has the same height at all the folded positions (Figure 4.4), it can damage a target that needed no correction (Figure 4.6), and the level it leaves is predicted by the geometry (Figure 4.5): it is a property of the order of the operations, not of the implementation. Applying SATA after the reconstruction does not help either, since one phase cannot undo a sum that has already been formed.
2. **The joint reconstruction solves both problems together.** Separating the parcels needs the phases corrected, and correcting each parcel needs the parcels separated. The joint reconstruction writes the phase of every ground position *inside* the model of the channels (Figure 5.4) and inverts model and topography together. With a flat DEM it is the standard reconstruction; on the periodic terrains where SATA is exact it gives the same result as SATA; everywhere else it keeps working.
3. **The whole-line version removes the last approximation.** The joint reconstruction in the sub-apertures still reads the DEM once per bin of a 40 ms block, i.e. it takes the terrain phase as constant over about ±600 m, and fails on steep terrain with long baselines (median −30.9 dB in the audit). Over the whole line the residual of every scatterer is a pointwise phase on the focused line, with no cell at all (median −41.4 dB).
4. **It uses only what a real case has.** The method receives the channels, the system geometry and the DEM; the positions of the targets are never used, and the unexplained part of the channels gives a check of the DEM.
5. **It degrades gracefully.** Where the baselines do not separate the parcels the regularisation leaves the ambiguity instead of amplifying noise, and the cost (a few seconds per line in the 1-D simulator) is that of a few tens of FFTs per channel and iteration.

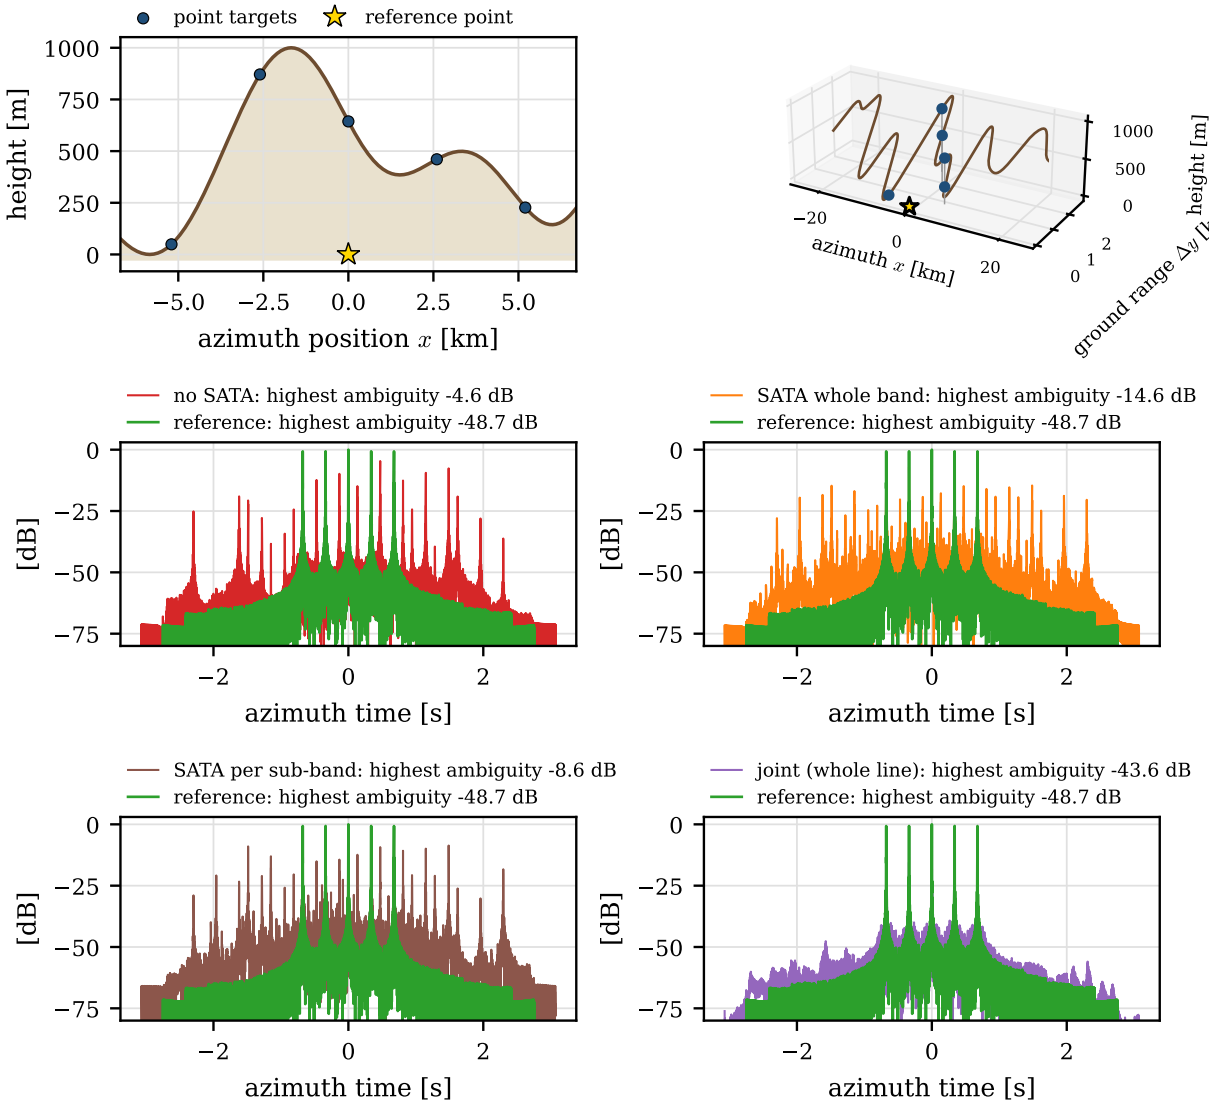


Figure 7.2: Five targets 2.6 km apart on random mountains (relief 1 000 m, correlation length 3 km), four channels, $b_{\text{xt}} = [0, 300, 119.8, 100.5]$ m. Top: the targets and the terrain near them (left) and along the whole azimuth line (right). Below: impulse responses of the reference and of no SATA, SATA whole band, SATA per sub-band and the joint reconstruction over the whole line. Every method reads only the DEM.

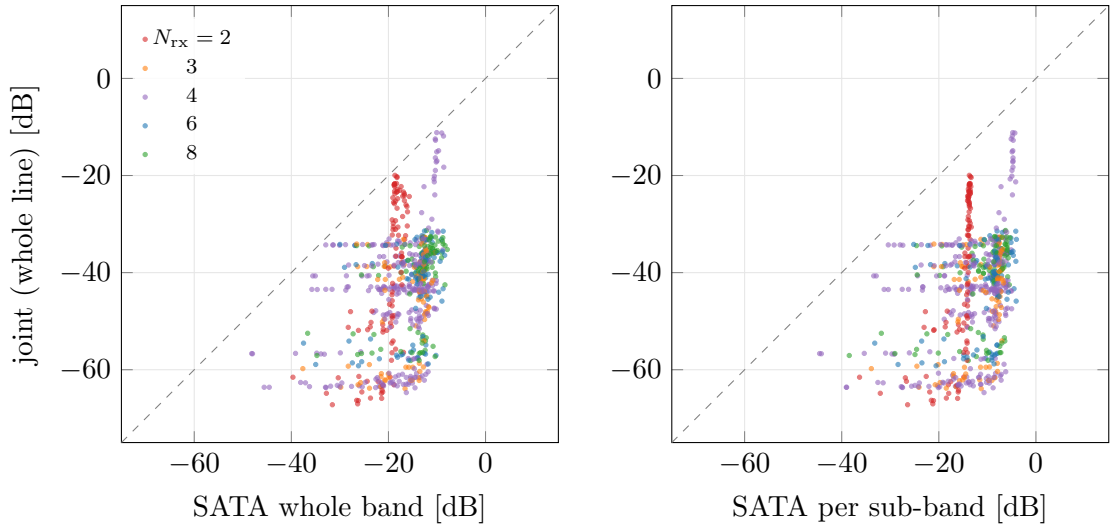


Figure 7.3: Highest ambiguity of the joint reconstruction over the whole line against SATA whole band (left) and SATA per sub-band (right), for the 680 runs of the audit (every method with the DEM only), coloured by the number of channels. Every point is below the diagonal.

7.3 Sweep of the final joint reconstruction in 1-D

The complete 1-D sweep runs the reference, the uncorrected reconstruction and the joint reconstruction over the whole line on 52 scenes × 20 baseline sets = 1040 runs. The medians are −47.6 dB (reference), −8.4 dB (no correction) and −40.6 dB (joint); 78 % of the runs are within 10 dB of the reference. Figure 7.4 splits the result by the number of channels and Figure 7.5 shows every run.

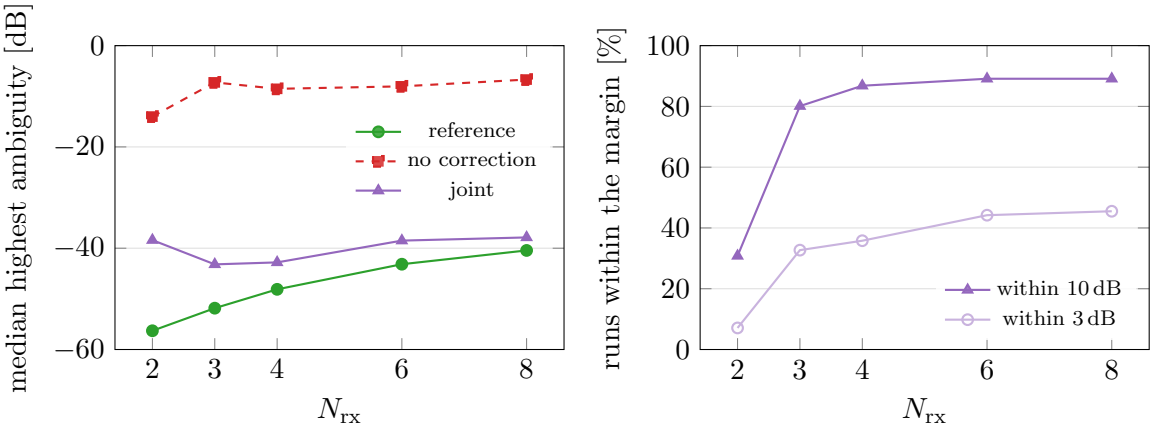


Figure 7.4: 1-D sweep by number of channels. Left: medians of the highest ambiguity. Right: fraction of runs in which the joint reconstruction is within 10 dB and 3 dB of the reference. The reference level rises with N_{rx} because PRF_{op} falls and more ambiguity orders fall inside the band.

The 83 runs that are not resolved (above −30 dB and more than 10 dB from the reference) are almost all configurations whose baselines cannot separate the parcels: 53 have two channels and 25 the equally spaced baselines $[0, 50, 100, 150]$ m; only 5 have four random baselines (3 with ${b_{xt}}^{max} = 300$ m, 2 with 500 m). Figure 7.7 shows a typical resolved case on a fractal alpine relief. Figure 7.6 shows the role of the number of channels on one scene with 49 targets on random mountains and ${b_{xt}}^{max} = 300$ m: with three or more channels the joint reconstruction is within

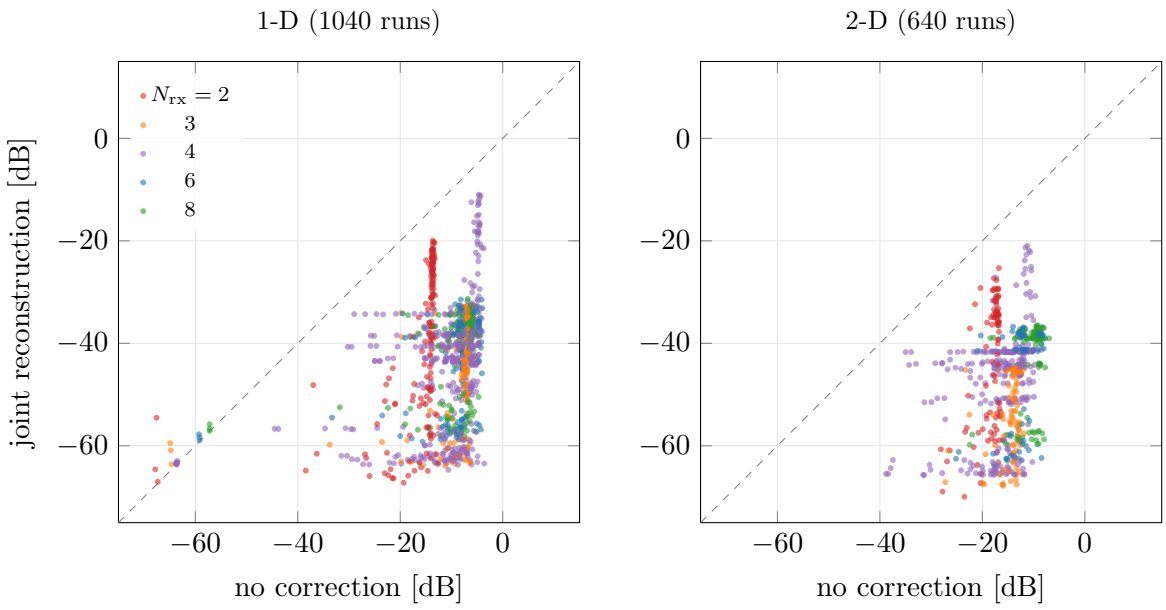


Figure 7.5: Every run of the 1-D and 2-D sweeps: highest ambiguity of the joint reconstruction against the uncorrected reconstruction, coloured by the number of channels. Points below the diagonal are improvements.

1 dB to 4 dB of the reference; with two channels there is no redundancy and it stops at −22.3 dB (reference −38.5 dB), still 8 dB better than no SATA.

7.4 Sweep of the joint reconstruction in 2-D

The 2-D sweep uses 32 3-D scenes (ramps along azimuth, range and oblique; sines and cosines; realistic and fractal terrains without layover; one target on a flank, five on a diagonal or a 7×7 grid) with the same 20 baseline sets, 640 runs in total. The medians are −47.6 dB (reference), −13.9 dB (no correction) and −44.4 dB (joint), all in the common range band. The joint reconstruction is never worse than the uncorrected one.

Table 7.2: 2-D sweep by number of channels: median gap between the joint reconstruction and the reference, and fraction of runs within 3 dB of the reference.

N_{rx}	runs	reference [dB]	no corr. [dB]	joint [dB]	joint − ref. [dB]	within 3 dB [%]
2	96	−60.8	−17.4	−39.3	22.9	6
3	96	−55.5	−13.8	−53.9	1.1	65
4	256	−45.2	−15.1	−44.6	0.3	84
6	96	−41.5	−12.2	−41.4	0.2	91
8	96	−39.1	−9.5	−39.1	0.2	96

With three or more channels the joint reconstruction is within a fraction of a dB of the reference in the median (Table 7.2). Of the 24 unresolved runs, 10 have two channels and 9 the equally spaced baselines $[0, 50, 100, 150]$ m; the others are 7×7 grids on a mountain, a valley or a volcano with long baselines.

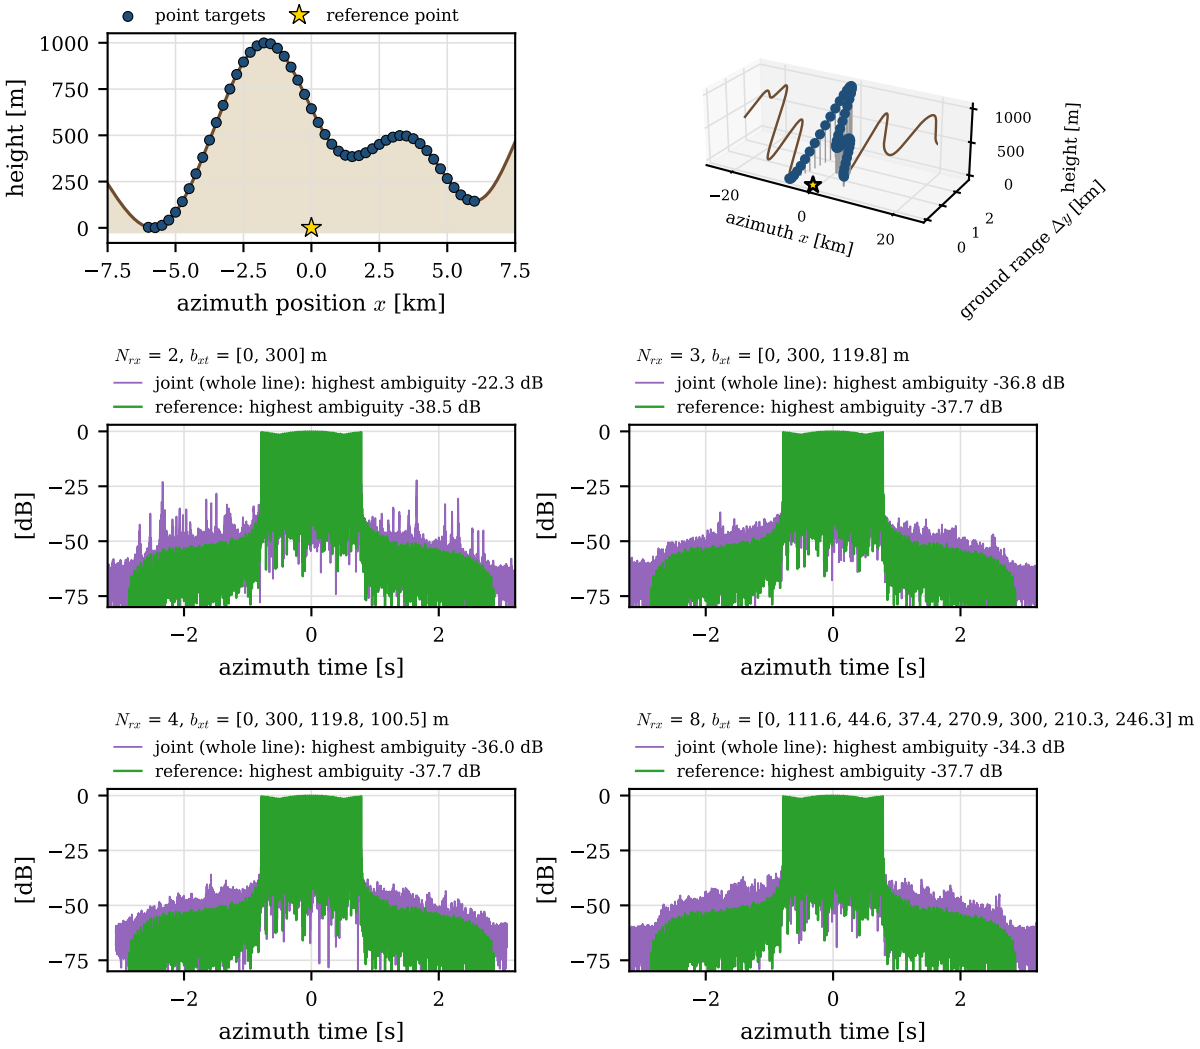


Figure 7.6: 49 targets 250 m apart on random mountains (relief 1 000 m), largest cross-track baseline 300 m, channel 0 on the transmitter. Top: the targets and the terrain near them (left) and along the whole azimuth line (right). Below: impulse responses of the reference and of the joint reconstruction over the whole line for $N_{\mathrm{rx}} = 2$, 3, 4 and 8.

7.5 Limits

DEM errors. The phase error is $\approx 9.0 \times 10^{-5}\, b_{\mathrm{xt}}\, \delta h$ rad, so the tolerance falls with the baseline (Figure 7.8). With an rms height error of 5 m, the joint reconstruction stays at −36 dB for $b_{\mathrm{xt}} = 135$ m but rises to −20 dB for 300 m.

To do: Confirm the scene and the error model behind Figure 7.8 (taken from the audit notes) and add the error-free point for ${b_{\mathrm{xt}}}^{\max} = 300$ m.

Noise. Relative to the standard reconstruction, the joint reconstruction raises the output noise by 0.3 dB to 4 dB in normal configurations, by 12 dB with equally spaced baselines and by 8 dB with two channels and 300 m. The regularisation keeps these numbers bounded where the problem is ill conditioned.

Configurations without redundancy. With two channels, or with baselines whose phases coincide at some heights, the data do not contain the information to separate the parcels. This is a property of the formation and should be considered in its design.

Not yet tested. Real antenna patterns (the model uses the simulated reference chirp), real orbits and baselines varying along the orbit, layover, distributed scenes in 2-D, noise combined with DEM errors, and real data.

7.6 Real terrain: SRTM lines over the Sierra Nevada

The scenes of the previous sections are synthetic. To check the joint reconstruction over the whole line on real relief, azimuth lines were cut from the SRTM 3″ DEM (about 90 m, WGS 84) [9] over the Sierra Nevada (California), using the file `raster/srtm90.tif` of `github.com/opengeos/data`. Every result is still one azimuth line (1-D): a straight 47 km profile of the DEM, sampled every 90 m and interpolated linearly between samples. The heights are taken relative to the mean height of the line, which is the reference height of the processing. Point targets of unit amplitude are placed on the terrain, the four channels are simulated with $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m (channel 0 on the transmitter), and the joint reconstruction receives the same SRTM profile as its DEM, never the targets. The metric is the highest ambiguity of Section 3.3; the reference is the monostatic full-rate line of the same scene. The script is `sata1d_code/joint_cases/real_dem/run_real_dem.py` and the figures are made by `real_dem_figures.py` in the same folder.

Table 7.3: The five named SRTM lines with five targets each (at $x = -5.0, -2.4, 0, 2.6, 5.3$ km, chosen so that no ghost at $\pm X$ or $\pm 2X$ falls on another target). Highest ambiguity in dB; the reference is −41.6 dB on every line (sidelobes of the neighbouring targets inside the ghost windows).

Line	Direction	Terrain height [m]	No correction [dB]	Joint [dB]
1 Yosemite Valley	north–south	1213–2893	−5.7	−41.5
2 Eastern Sierra escarpment	west–east	1198–4053	−8.2	−41.5
3 High Sierra, Tuolumne	north–south	2351–3635	−6.0	−41.6
4 Western foothills	west–east	52–932	−6.1	−41.4
5 Mono basin	north–south	1948–3432	−6.2	−41.0

Figures 7.10 to 7.12 show three of the lines. Without correction every target leaves ghosts at $\pm X$ and $\pm 2X$ ($X = 6.23$ km) between −8.2 dB and −5.7 dB; the joint reconstruction brings all

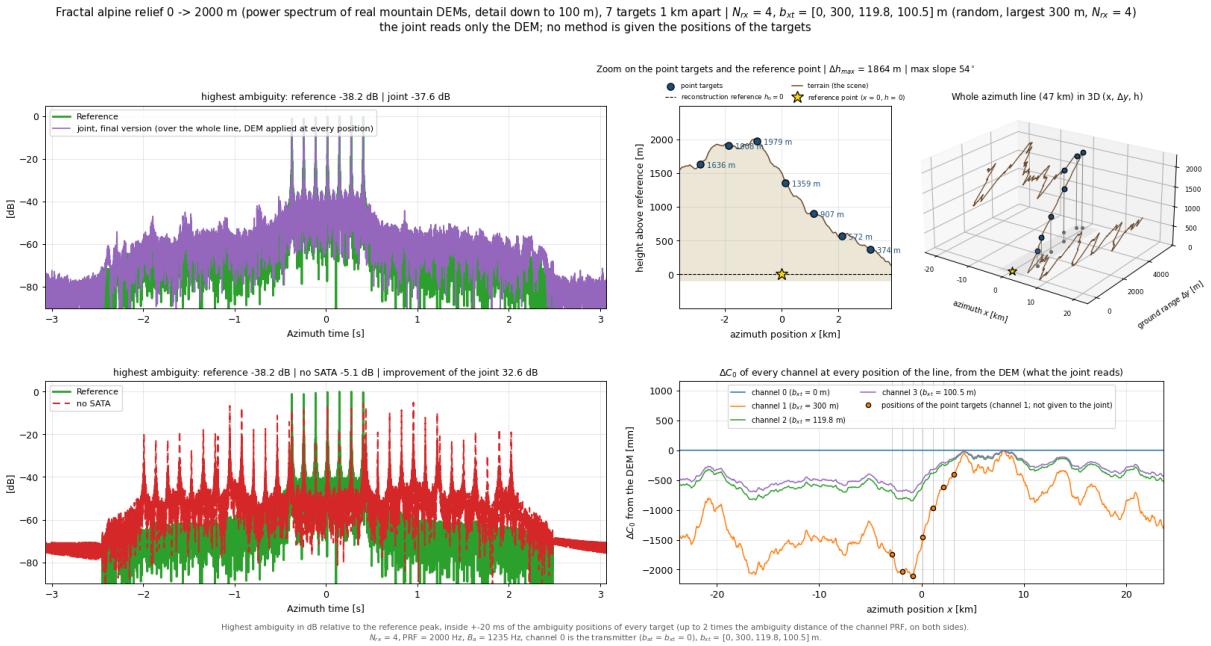


Figure 7.7: A page of the 1-D sweep: seven targets on a fractal alpine relief (0 m to 2 000 m, slopes up to 54°), $N_{\mathrm{rx}} = 4$, $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m. The joint reconstruction (top left) reaches the reference level (−37.6 dB against −38.2 dB); without correction the highest ambiguity is −5.1 dB. Bottom right: the ΔC_0 of every channel along the whole line, read from the DEM.

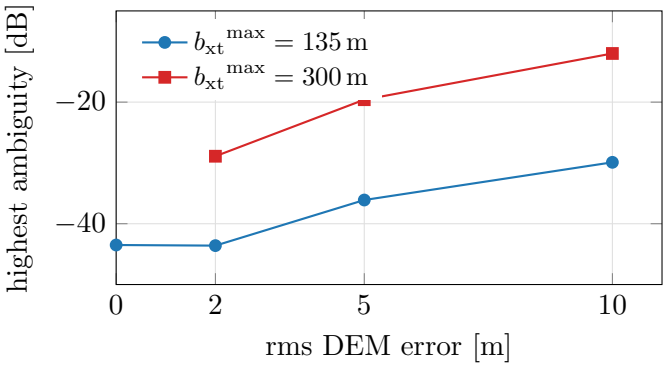


Figure 7.8: Joint reconstruction over the whole line with a random DEM error (five targets on the mountain of Figure 5.8).

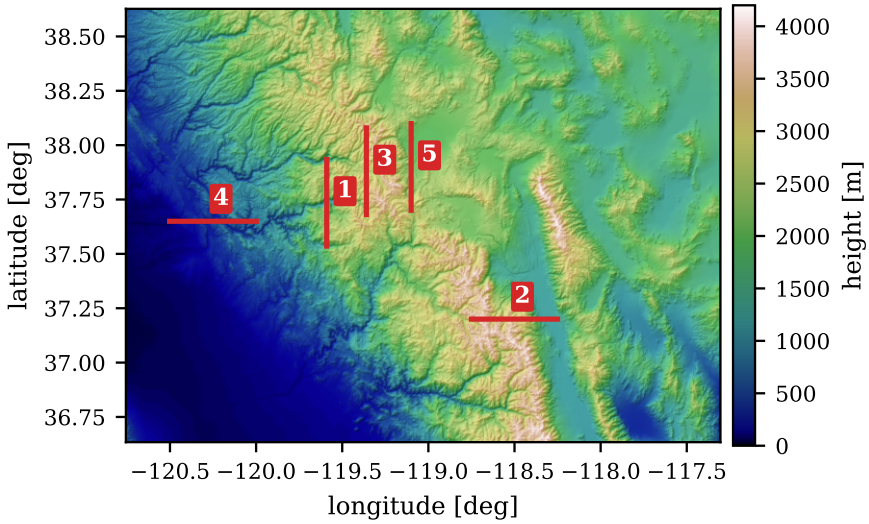


Figure 7.9: SRTM 3″ DEM of the Sierra Nevada with the five named lines of Table 7.3 (red, 47 km each). The map only shows where the lines come from: every result is computed on one line.

of them to the level of the reference, within 0.6 dB on the five lines (Table 7.3). The Yosemite line has steep valley walls (about 1 700 m of relief over a few kilometres), the escarpment line goes from the Sierra crest down to the floor of Owens Valley (2 800 m), and the High Sierra line has many short slopes, so that the height, and with it the phase of every channel, changes every few hundred metres. None of these is a difficulty for the joint reconstruction, because it applies the DEM phase at every position of the line.

Figure 7.13 zooms on the Yosemite target that is farthest from the reference height (on the valley floor, 487 m below it). Without correction the peak is 4.4 dB low and 29.7° off, because part of its energy went to the ghosts; the ghost at $+X$ is at −10.2 dB. With the joint reconstruction the peak equals the reference within 0.01 dB and 0.02°, and the ghost drops to −43.1 dB, below the sidelobes of the reference around it.

Many targets. With 20 targets every 750 m (from −7.1 km to 7.1 km) every ghost falls between two targets, and the highest ambiguity is the worst of 80 ghost windows per line. The reference reaches only −32.0 dB here, from the sidelobes of the neighbours inside the windows. Without correction the five lines give −5.7 dB to −5.1 dB; the joint reconstruction gives −31.5 dB to −28.1 dB, within 4 dB of the reference (Figure 7.14).

200 random lines. To avoid choosing the lines by hand, 200 lines were drawn at random over the whole DEM (random centre and direction, the full 47 km inside the raster), with the five targets of Table 7.3. The relief of a line reaches 2 930 m; half of the lines use $b_{\mathrm{xt}} = [0, 300, 119.8, 100.5]$ m and half $b_{\mathrm{xt}} = [0, 134.9, 20.5, 8.3]$ m (largest 135 m). Figure 7.15 gives the result: the median highest ambiguity goes from −8.3 dB without correction to −41.5 dB with the joint reconstruction (reference −41.6 dB), a median gain of 33 dB. In the 192 lines with more than 200 m of relief the gain is at least 28 dB in 92 % of them, and the joint reconstruction is within 10 dB of the reference in 199 of the 200 lines. It is never worse than no correction by more than 0.03 dB; that happens only on a flat line (16 m of relief), where both are already at the reference. The gain grows with the relief, because the uncorrected ghosts grow with it while the joint result stays at the reference.

These runs share the limits of the rest of the chapter: the DEM given to the reconstruction

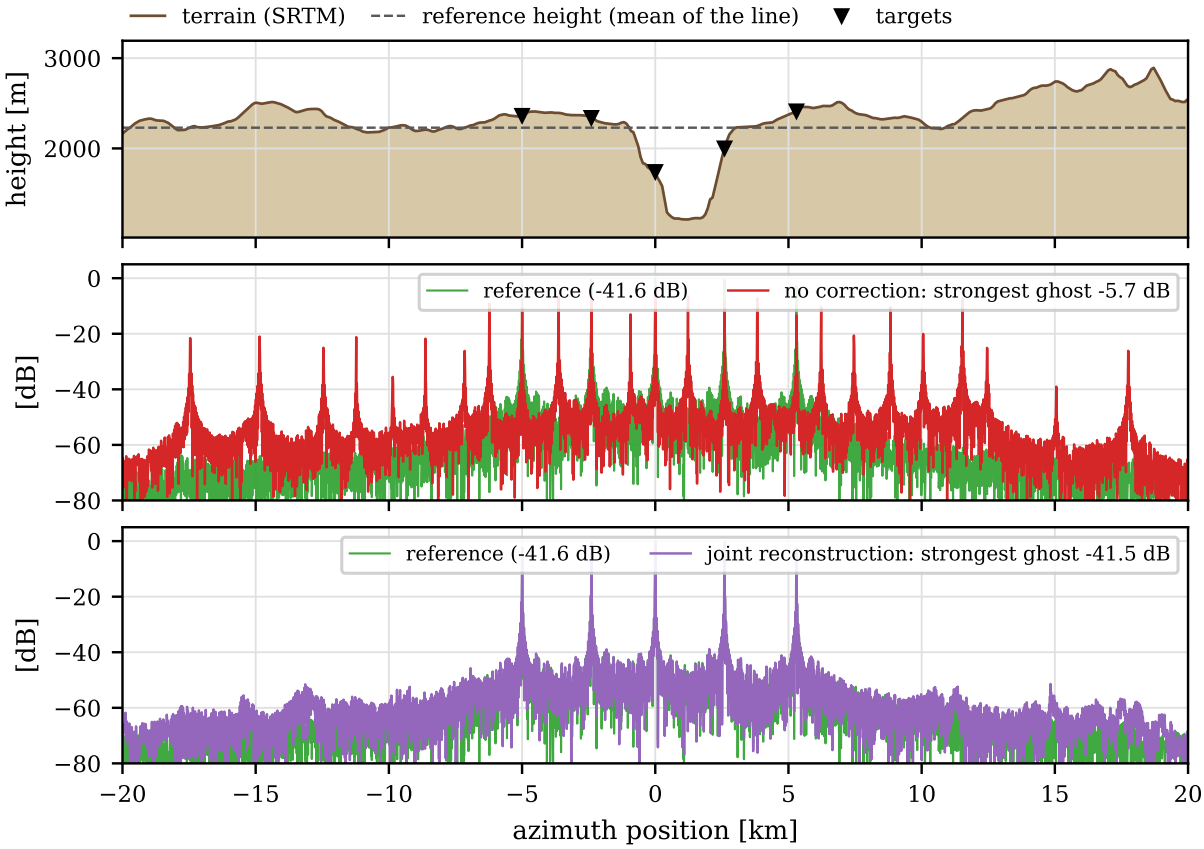


Figure 7.10: Yosemite Valley (north–south, line 1). Top: SRTM profile, reference height and the five targets. Middle and bottom: focused line in dB relative to the peak of the reference, without correction and with the joint reconstruction; the reference (green) is drawn below both.

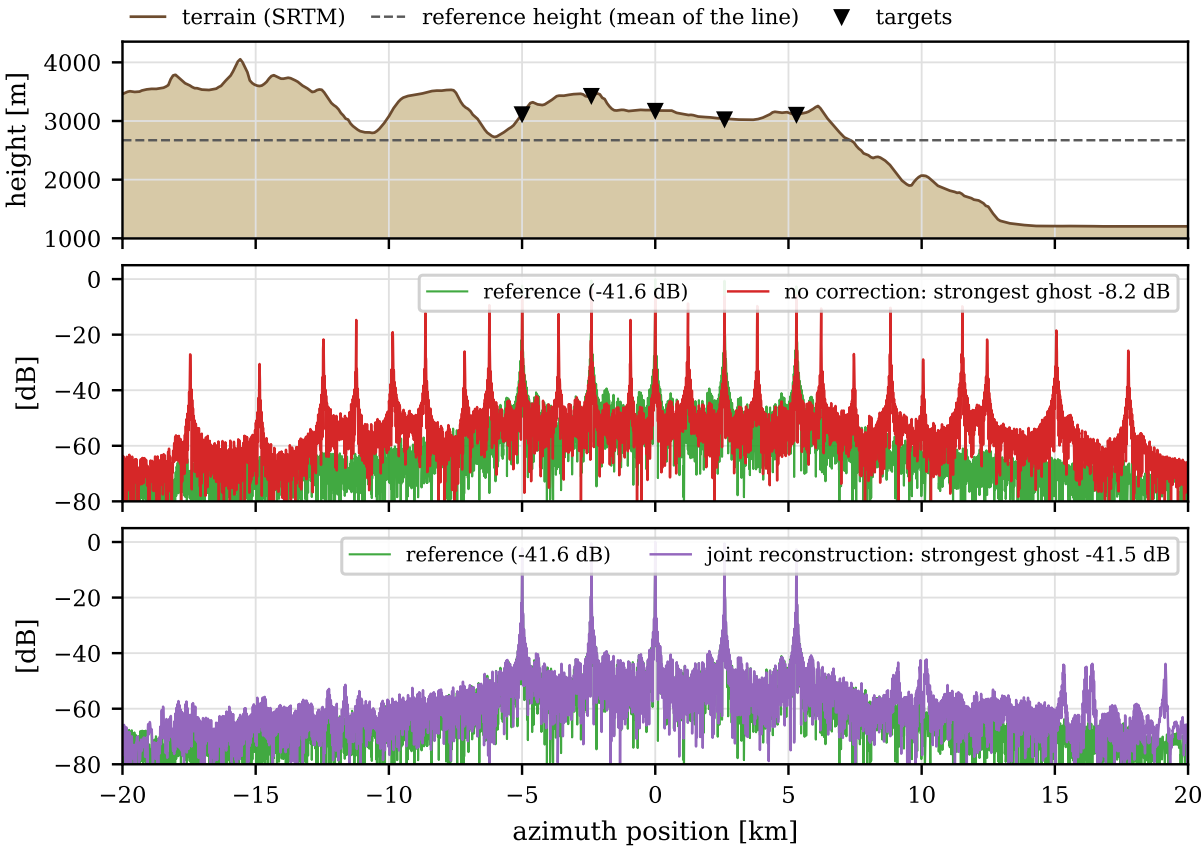


Figure 7.11: Eastern Sierra escarpment (west–east, line 2), from the crest to the floor of Owens Valley; same layout as Figure 7.10.

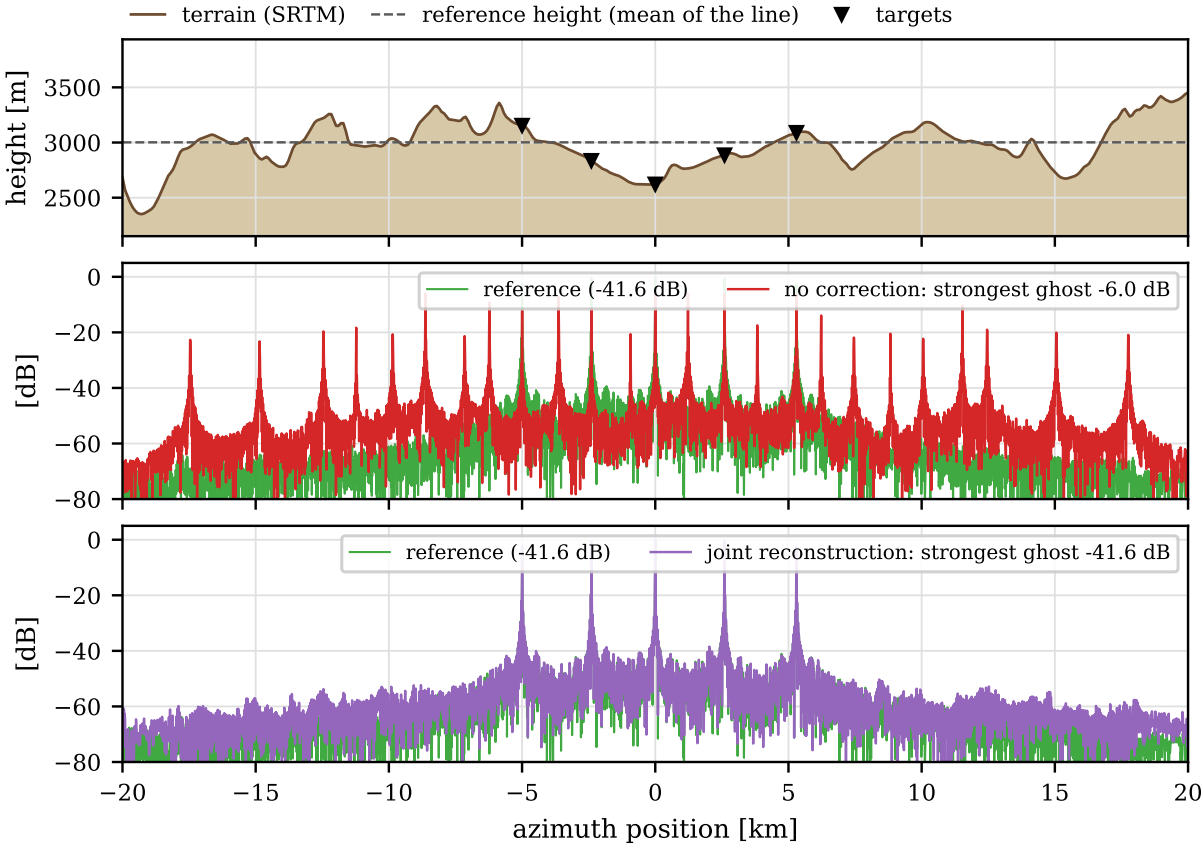


Figure 7.12: High Sierra near Tuolumne (north–south, line 3): rugged peaks and valleys; same layout as Figure 7.10.

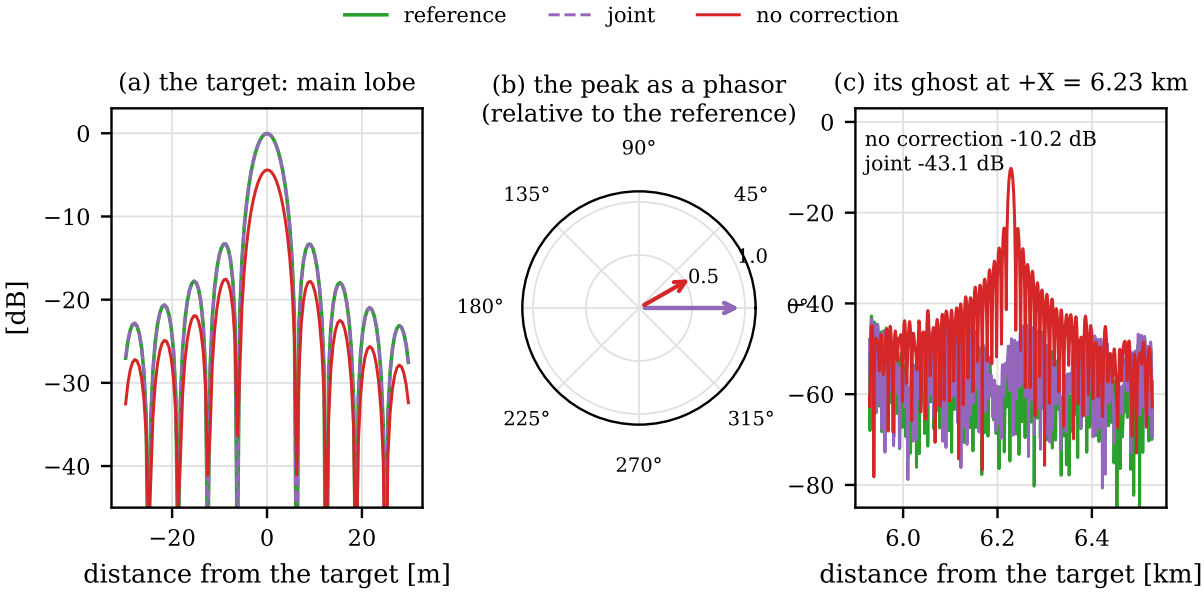


Figure 7.13: Yosemite, valley-floor target 487 m below the reference height. (a) Main lobe; the joint line lies on the reference. (b) Complex peak relative to the reference. (c) The window of its ghost at $+X$. All levels in the processed band B_a, relative to the peak of the reference.

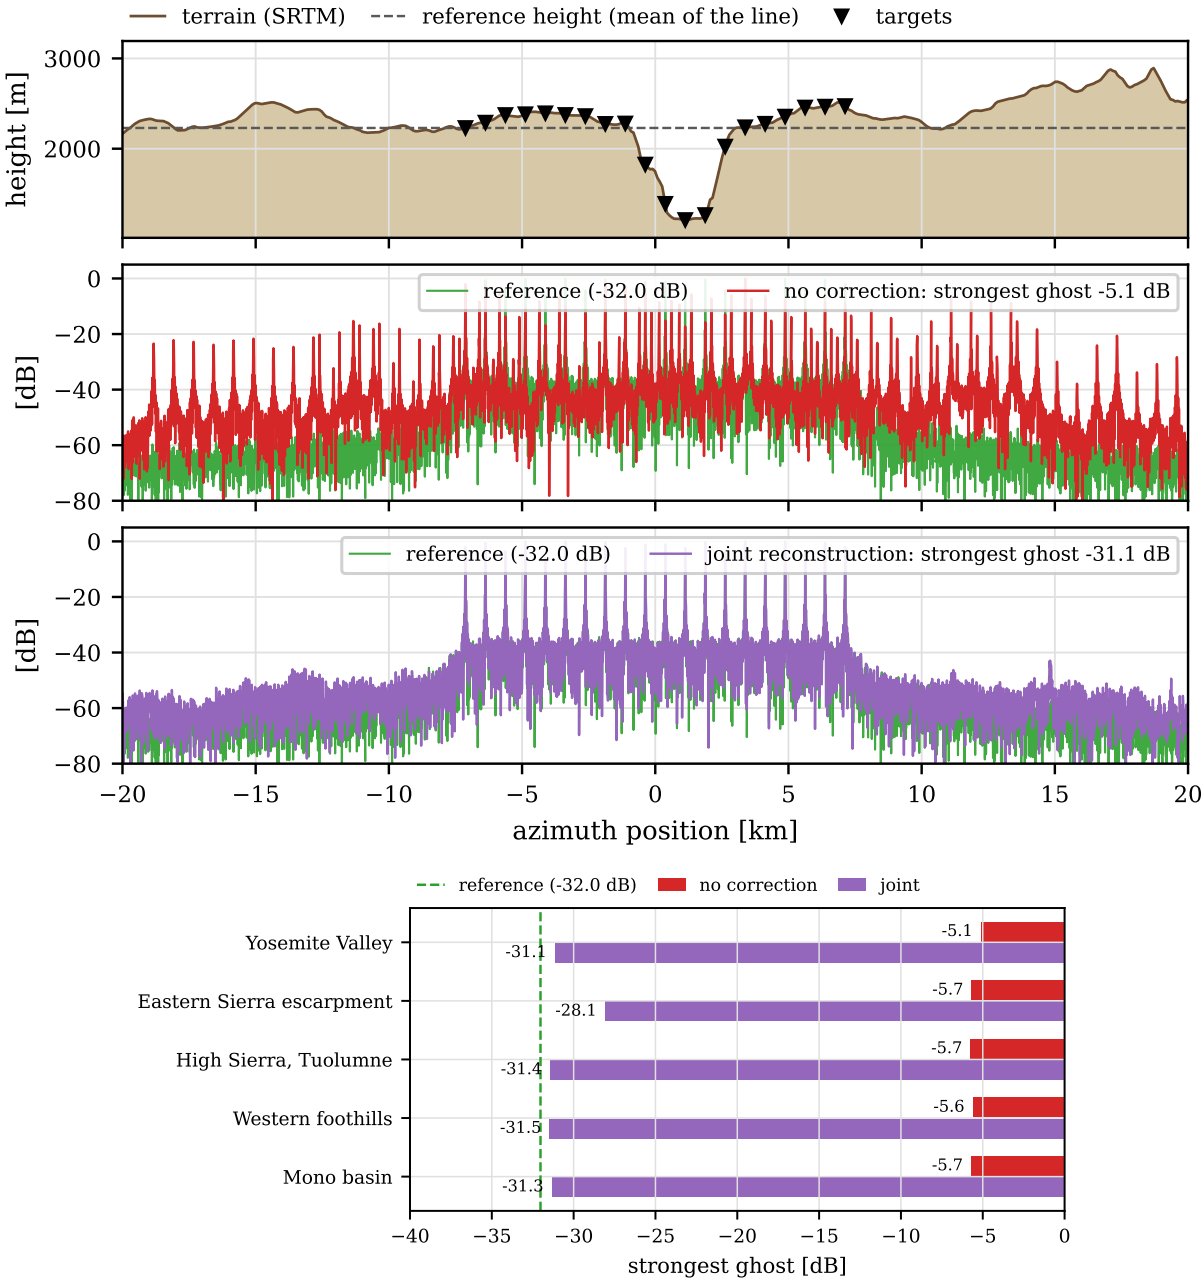


Figure 7.14: Top: Yosemite line with 20 targets every 750 m; same layout as Figure 7.10. Bottom: highest ambiguity of the five lines with 20 targets each.

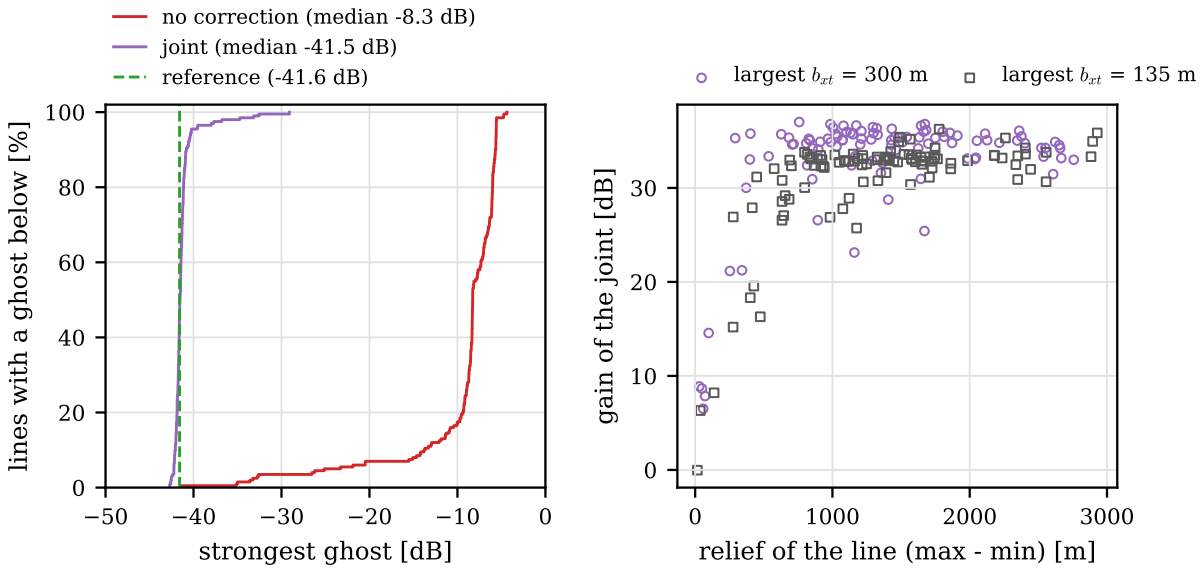


Figure 7.15: 200 random SRTM lines over the Sierra Nevada, five targets each. Left: cumulative distribution of the highest ambiguity. Right: gain of the joint reconstruction over no correction against the relief of the line, for the two baseline sets.

is exactly the terrain of the simulation (with DEM errors the floor rises as in Figure 7.8), the antenna pattern is the simulated one, and each line is processed in 1-D.

8 Conclusions and outlook

8.1 Conclusions

- Cross-track baselines turn topography into a channel-dependent phase that the standard multichannel reconstruction cannot absorb, $\varphi_i \approx 2\pi b_{\mathrm{xt},i} h/(\lambda r_0 \tan\theta)$. Without correction the median highest ambiguity of the sweeps is $-8.4\,\mathrm{dB}$ in 1-D and $-13.9\,\mathrm{dB}$ in 2-D, against a reference near $-48\,\mathrm{dB}$.
- SATA computes the right phase from the DEM, but applied to each aliased channel before the reconstruction it can give only one phase to the parcels of positions $X \approx 6.2\,\mathrm{km}$ apart that share a Doppler bin. It is exact only when the terrain has the same height at all those positions; elsewhere it leaves a structural residual that is predicted from the geometry within about 1 dB, and with a DEM only it improves the median by 5.5 dB at best.
- Placing the DEM phase inside the inversion solves this. The joint reconstruction over the whole line models every channel from the focused line and the DEM and inverts it by regularised least squares. Using only the system parameters and the DEM it reaches a median of $-40.6\,\mathrm{dB}$ in 1-D (78 % of the runs within 10 dB of the reference, about 90 % for four or more channels) and $-44.4\,\mathrm{dB}$ in 2-D, where with three or more channels it is within 1 dB of the reference in the median.
- On real terrain (200 azimuth lines cut from the SRTM DEM of the Sierra Nevada, relief up to 2 930 m) the joint reconstruction lowers the median highest ambiguity from $-8.3\,\mathrm{dB}$ to $-41.5\,\mathrm{dB}$, at the level of the reference ($-41.6\,\mathrm{dB}$), and is within 10 dB of it in 199 of the 200 lines (Section 7.6).
- The remaining failures are formations whose baselines do not separate the folded parcels (two channels, equally spaced baselines). The regularisation leaves the ambiguity there instead of amplifying noise.
- The accuracy is limited by the DEM: the phase error grows as $b_{\mathrm{xt}}\,\delta h$, so long baselines need accurate DEMs.

8.2 Outlook

- Validation with realistic antenna patterns, orbits and baselines that vary along the orbit, and eventually with real multistatic data.
- 2-D: distributed scenes, layover and shadow, finer range resolution (envelope interpolation) and wide swaths (range-dependent kernels in blocks).
- Joint estimation of a DEM correction, using the unexplained part of the channels as the cost.
- Use of the conditioning analysis (Section 3.2) as a design rule for the cross-track baselines of a formation.
- Computational cost: block-wise or GPU implementation of the model and its adjoint for full scenes.

Bibliography

[1] G. Krieger, N. Gebert, and A. Moreira, "Unambiguous SAR signal reconstruction from nonuniform displaced phase center sampling," *IEEE Geoscience and Remote Sensing Letters*, vol. 1, no. 4, pp. 260–264, 2004. (cited on pages 1 and 9)

[2] N. Gebert, G. Krieger, and A. Moreira, "Digital beamforming on receive: Techniques and optimization strategies for high-resolution wide-swath SAR imaging," *IEEE Transactions on Aerospace and Electronic Systems*, vol. 45, no. 2, pp. 564–592, 2009. (cited on pages 1 and 9)

[3] N. Sakar, M. Rodriguez-Cassola, P. Prats-Iraola, and A. Moreira, "Azimuth reconstruction algorithm for multistatic SAR formations with large along-track baselines," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 58, no. 3, pp. 1931–1940, 2020. (cited on pages 1 and 5)

[4] N. Sakar, *Processing Approaches for Multistatic Large Along-Track Baseline SAR Constellations*. PhD thesis, Technische Universität München, Munich, Germany, 2022. (cited on pages 1 and 5)

[5] P. Prats, A. Reigber, and J. J. Mallorqui, "Topography-dependent motion compensation for repeat-pass interferometric SAR systems," *IEEE Geoscience and Remote Sensing Letters*, vol. 2, no. 2, pp. 206–210, 2005. (cited on pages 1 and 13)

[6] P. Prats Iraola, *Airborne Differential SAR Interferometry*. PhD thesis, Universitat Politècnica de Catalunya, Barcelona, Spain, Nov. 2005. (cited on pages 1 and 13)

[7] M. R. Hestenes and E. Stiefel, "Methods of conjugate gradients for solving linear systems," *Journal of Research of the National Bureau of Standards*, vol. 49, no. 6, pp. 409–436, 1952. (cited on page 20)

[8] L. N. Trefethen and D. Bau, *Numerical Linear Algebra*. Philadelphia: SIAM, 1997. (cited on page 20)

[9] T. G. Farr, P. A. Rosen, E. Caro, R. Crippen, R. Duren, S. Hensley, M. Kobrick, M. Paller, E. Rodriguez, L. Roth, D. Seal, S. Shaffer, J. Shimada, J. Umland, M. Werner, M. Oskin, D. Burbank, and D. Alsdorf, "The Shuttle Radar Topography Mission," *Reviews of Geophysics*, vol. 45, no. 2, p. RG2004, 2007. (cited on page 40)